



Methodology of extraction of reliable energy data on a basic block level

Dimitrios Stamatios Bouras - el17072@mail.ntua.gr



Laboratory: Microprocessors and Digital systems Lab

University: National Technical University of Athens

Department: Electrical and Computer Engineering

Supervising Professor: Dimitrios Soudris

Supervising Phd candidate: Christos Lamprakos

Github repo: https://github.com/jimbou/dataset_tool.git

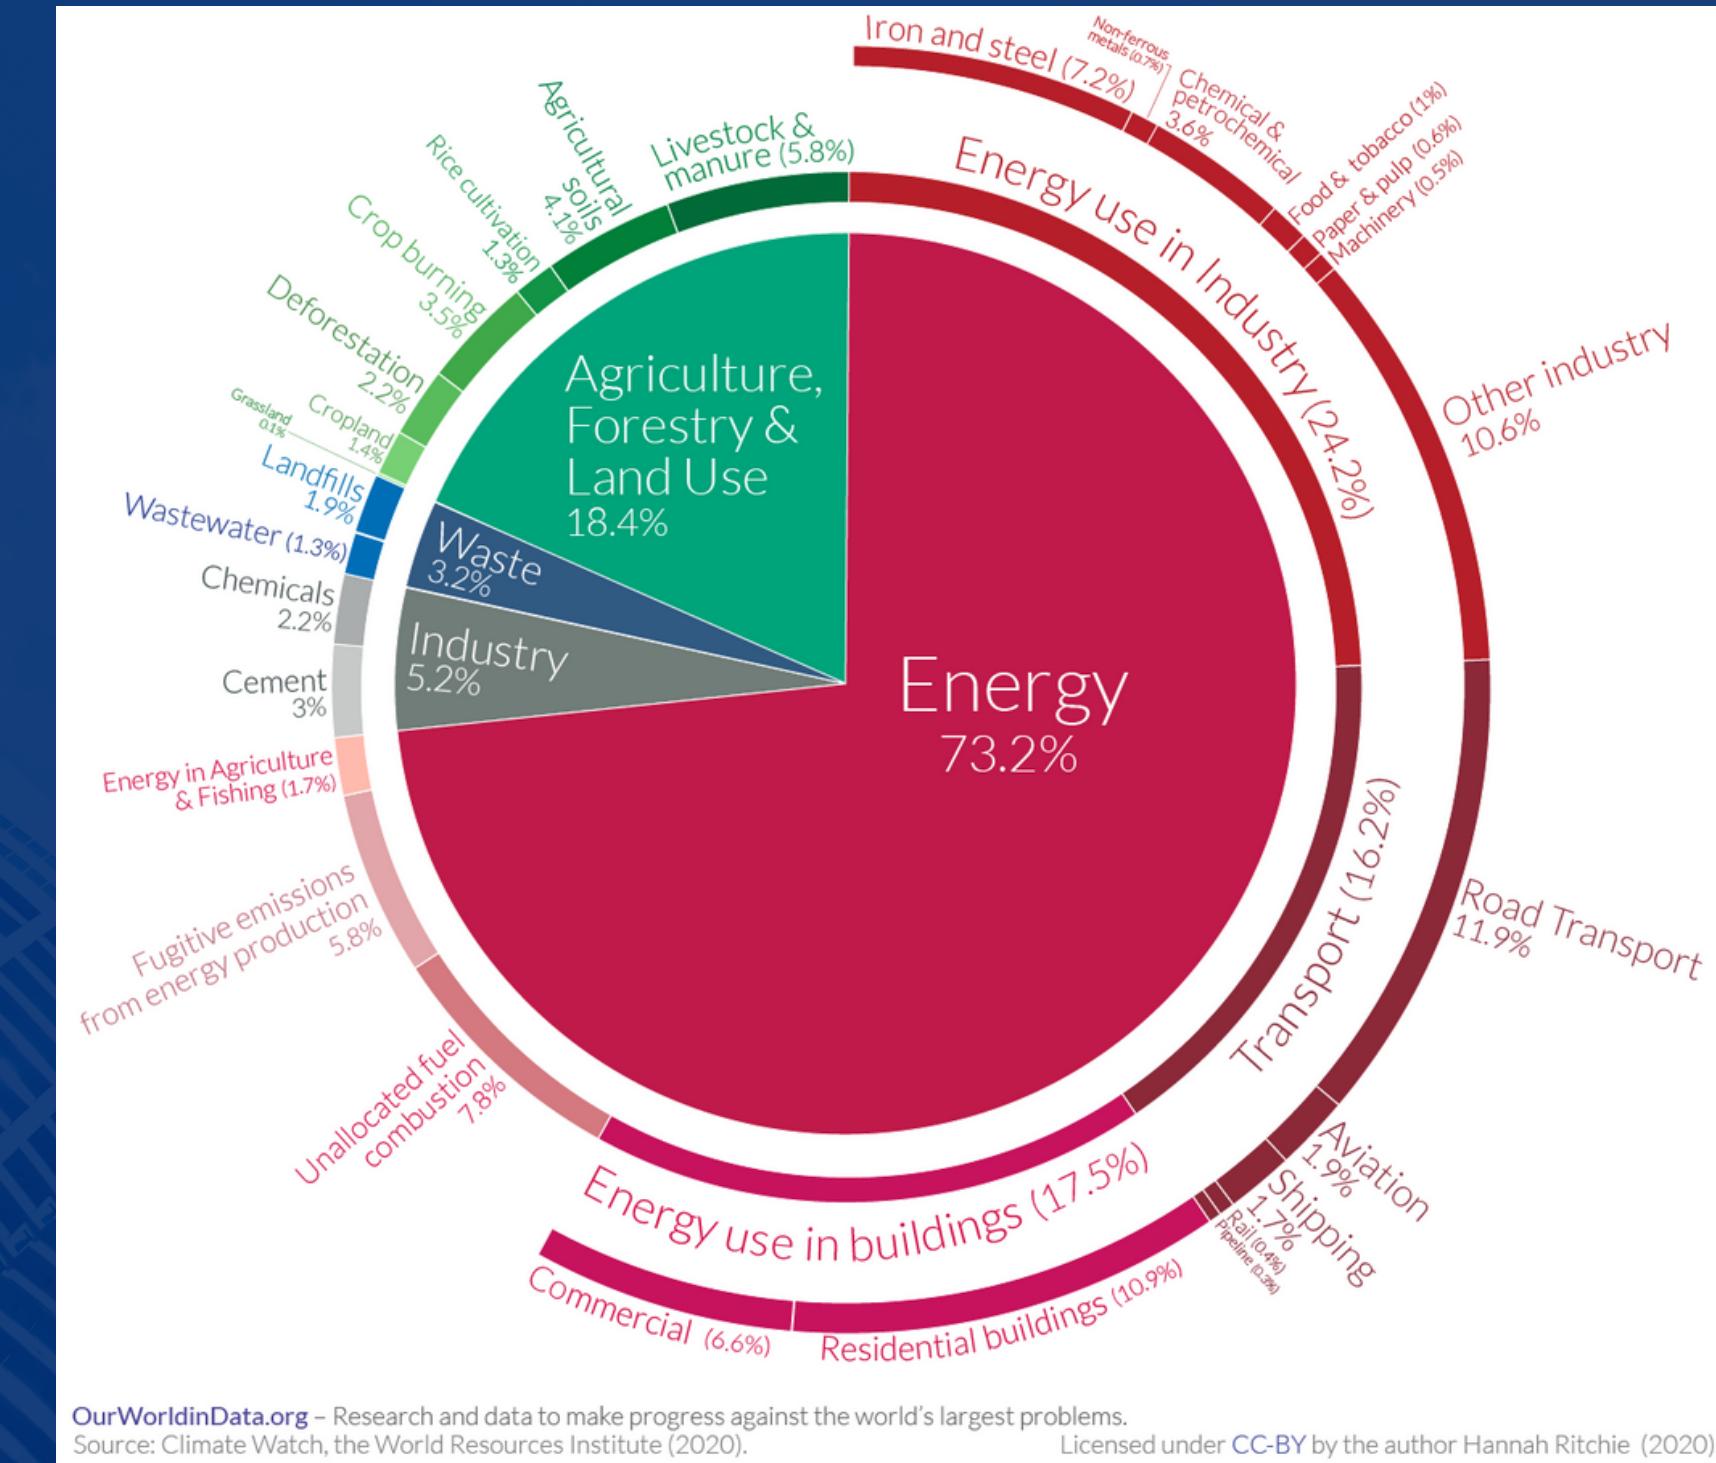


Energy efficient Software

- Reduced energy cost
- Reduced Heat dissipation
- Reduce energy footprint - GHG emmission
- Longer battery life



Greenhouse gas emissions per sector



ICT ~ 4%



Energy efficient Hardware

- Larger improvements with less change
- Improvements easily observed

but

- Energy cost per calculation decreasing but calculations' amount increasing
- Unable to satisfy increasing energy needs



To achieve



Energy efficient
software

First



Software energy
measuring tool



Why basic blocks?

- Higher Granularity
- Higher Accuracy
- More flexibility for developers
- Can't get any lower practically with energy reduction in mind

Basic block: a straight-line code sequence with one branch in and one branch out and no branches out.



Infrastructure

Computer specs :

Specifications of the computer utilized for this thesis	
Operating System	Linux Ubuntu 20.04.1
Core version	5.15.0-58 -generic
Architecture	x86 ₆ 4
Processor	Intel Core i7 – 6700
Processor frequency	3.4 GHz
Bridge	Xeon E3-1200 v5/E3-1500 v5/6th Gen
Bus	100 Series/C230 3.0 xHCI Controller
Communication	100 Series/C230 3.0 MEI Controller 1
RAM	32GiB System memory



Energy Measurements

Intel RAPL

- set of registers providing energy and power consumption information
- uses a software power model based on performance counters
- Refresh frequency 1KHz
- Available measurements for core , uncore, RAM and GPU
- High accuracy



LLVM

- LLVM is a set of compiler and toolchain technologies
- Designed around a language-independent intermediate representation (IR)
- language-independent instruction set and type system
- LLVM passes perform analysis, transformations and optimizations on the code



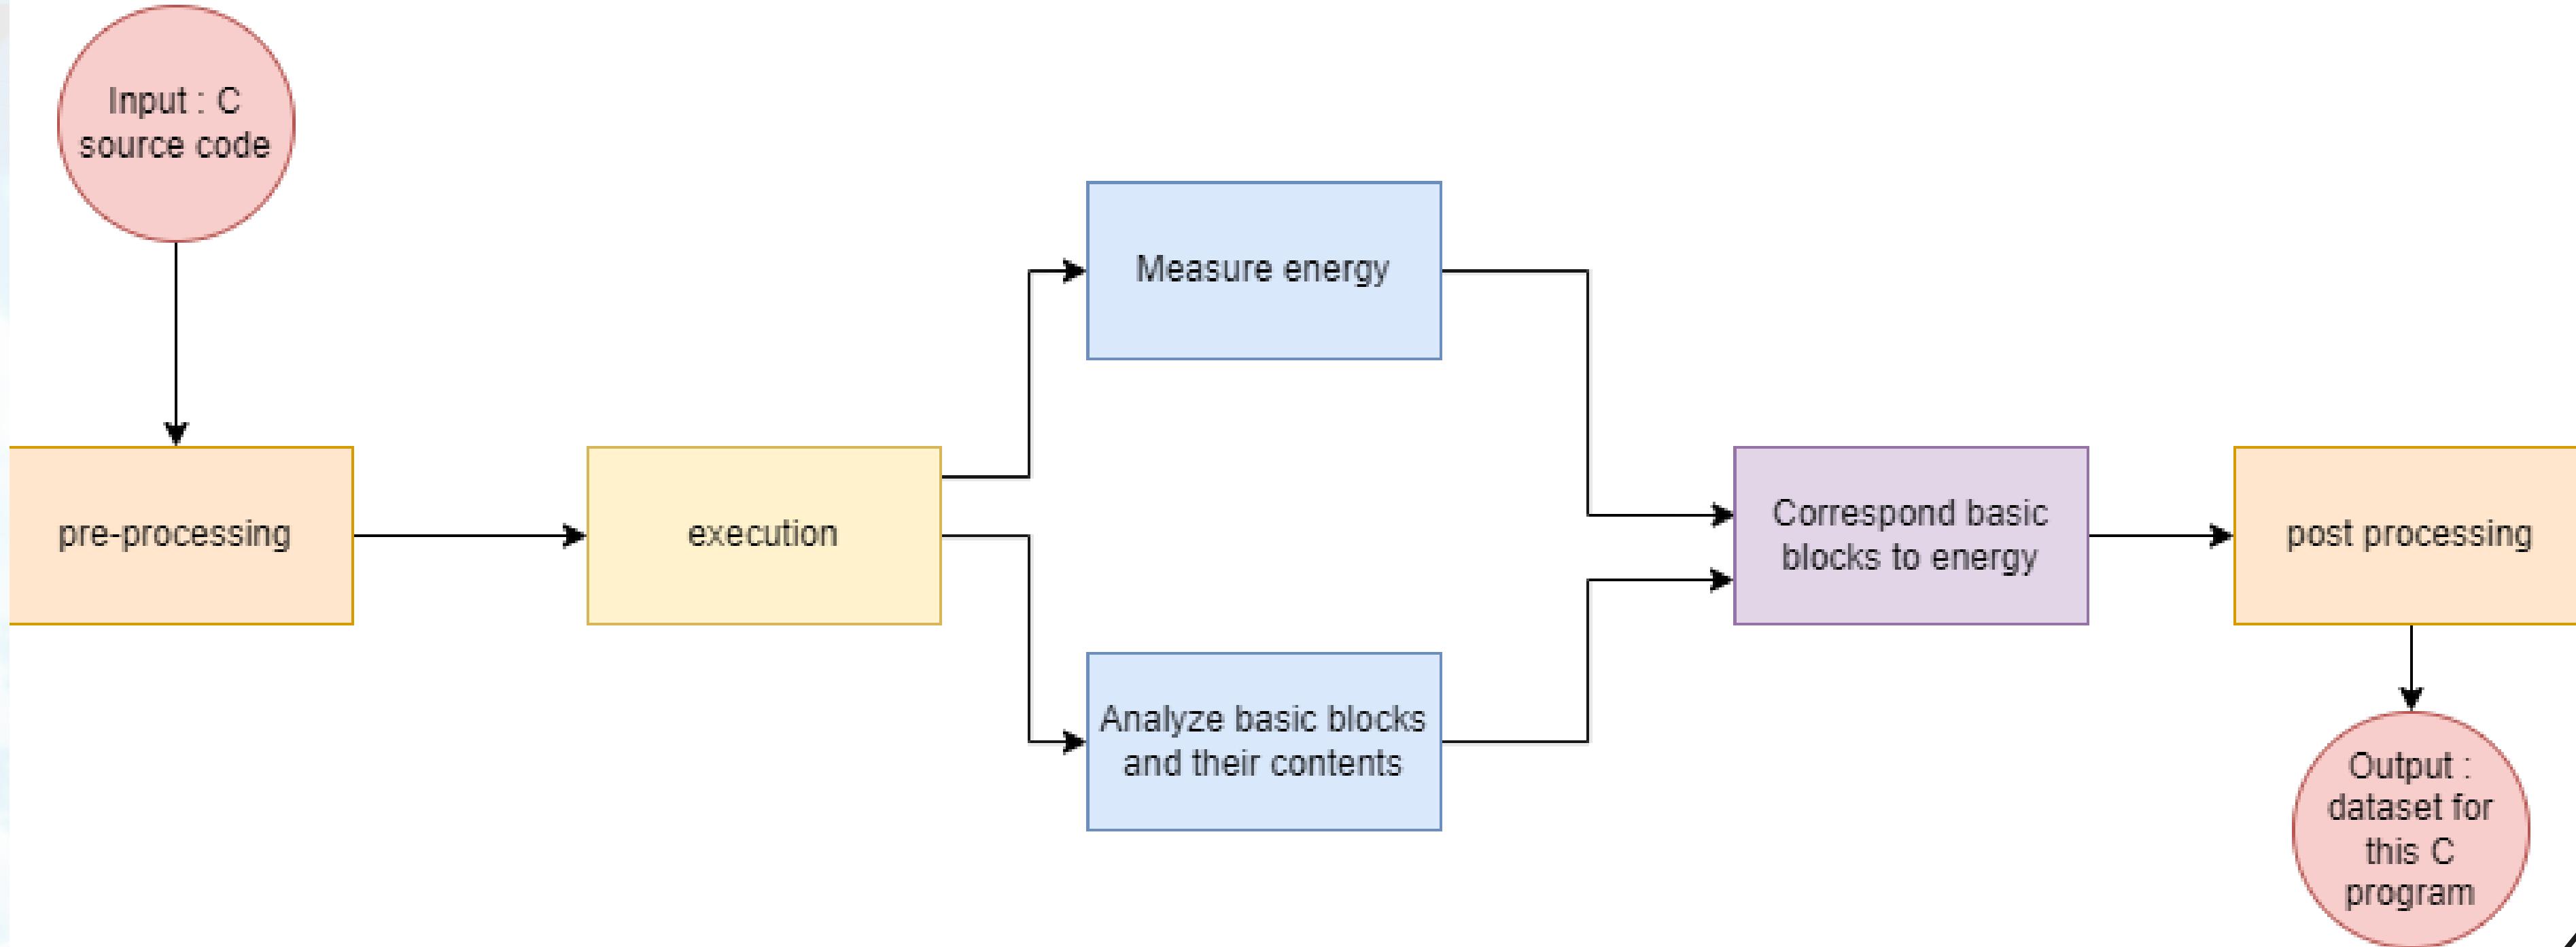


3 different approaches:

- Individual basic block measuring with RAPL calls
- Binary Lifting to LLVM IR
- Post analysis of execution trace



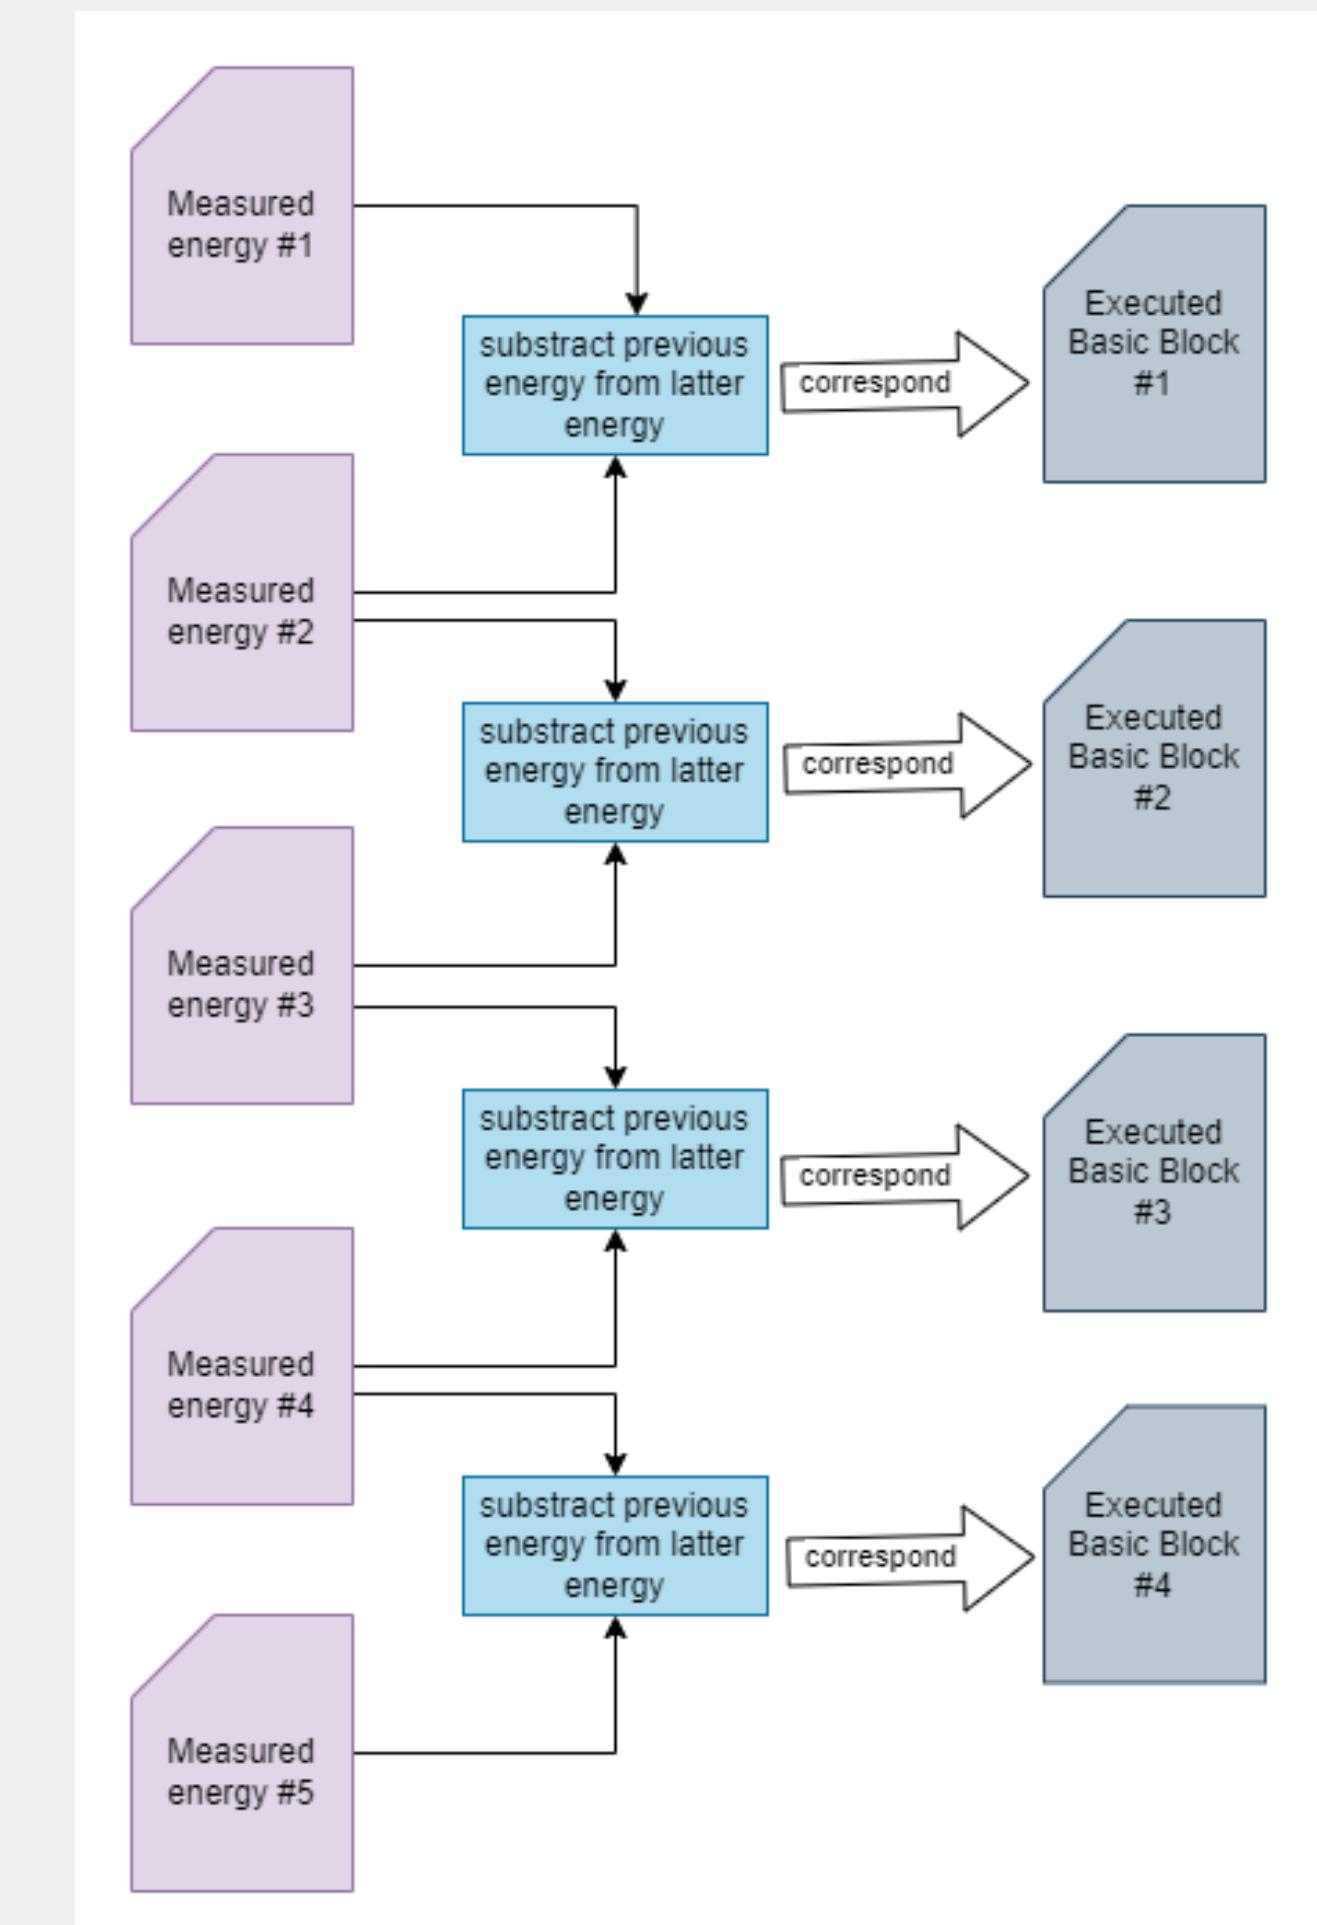
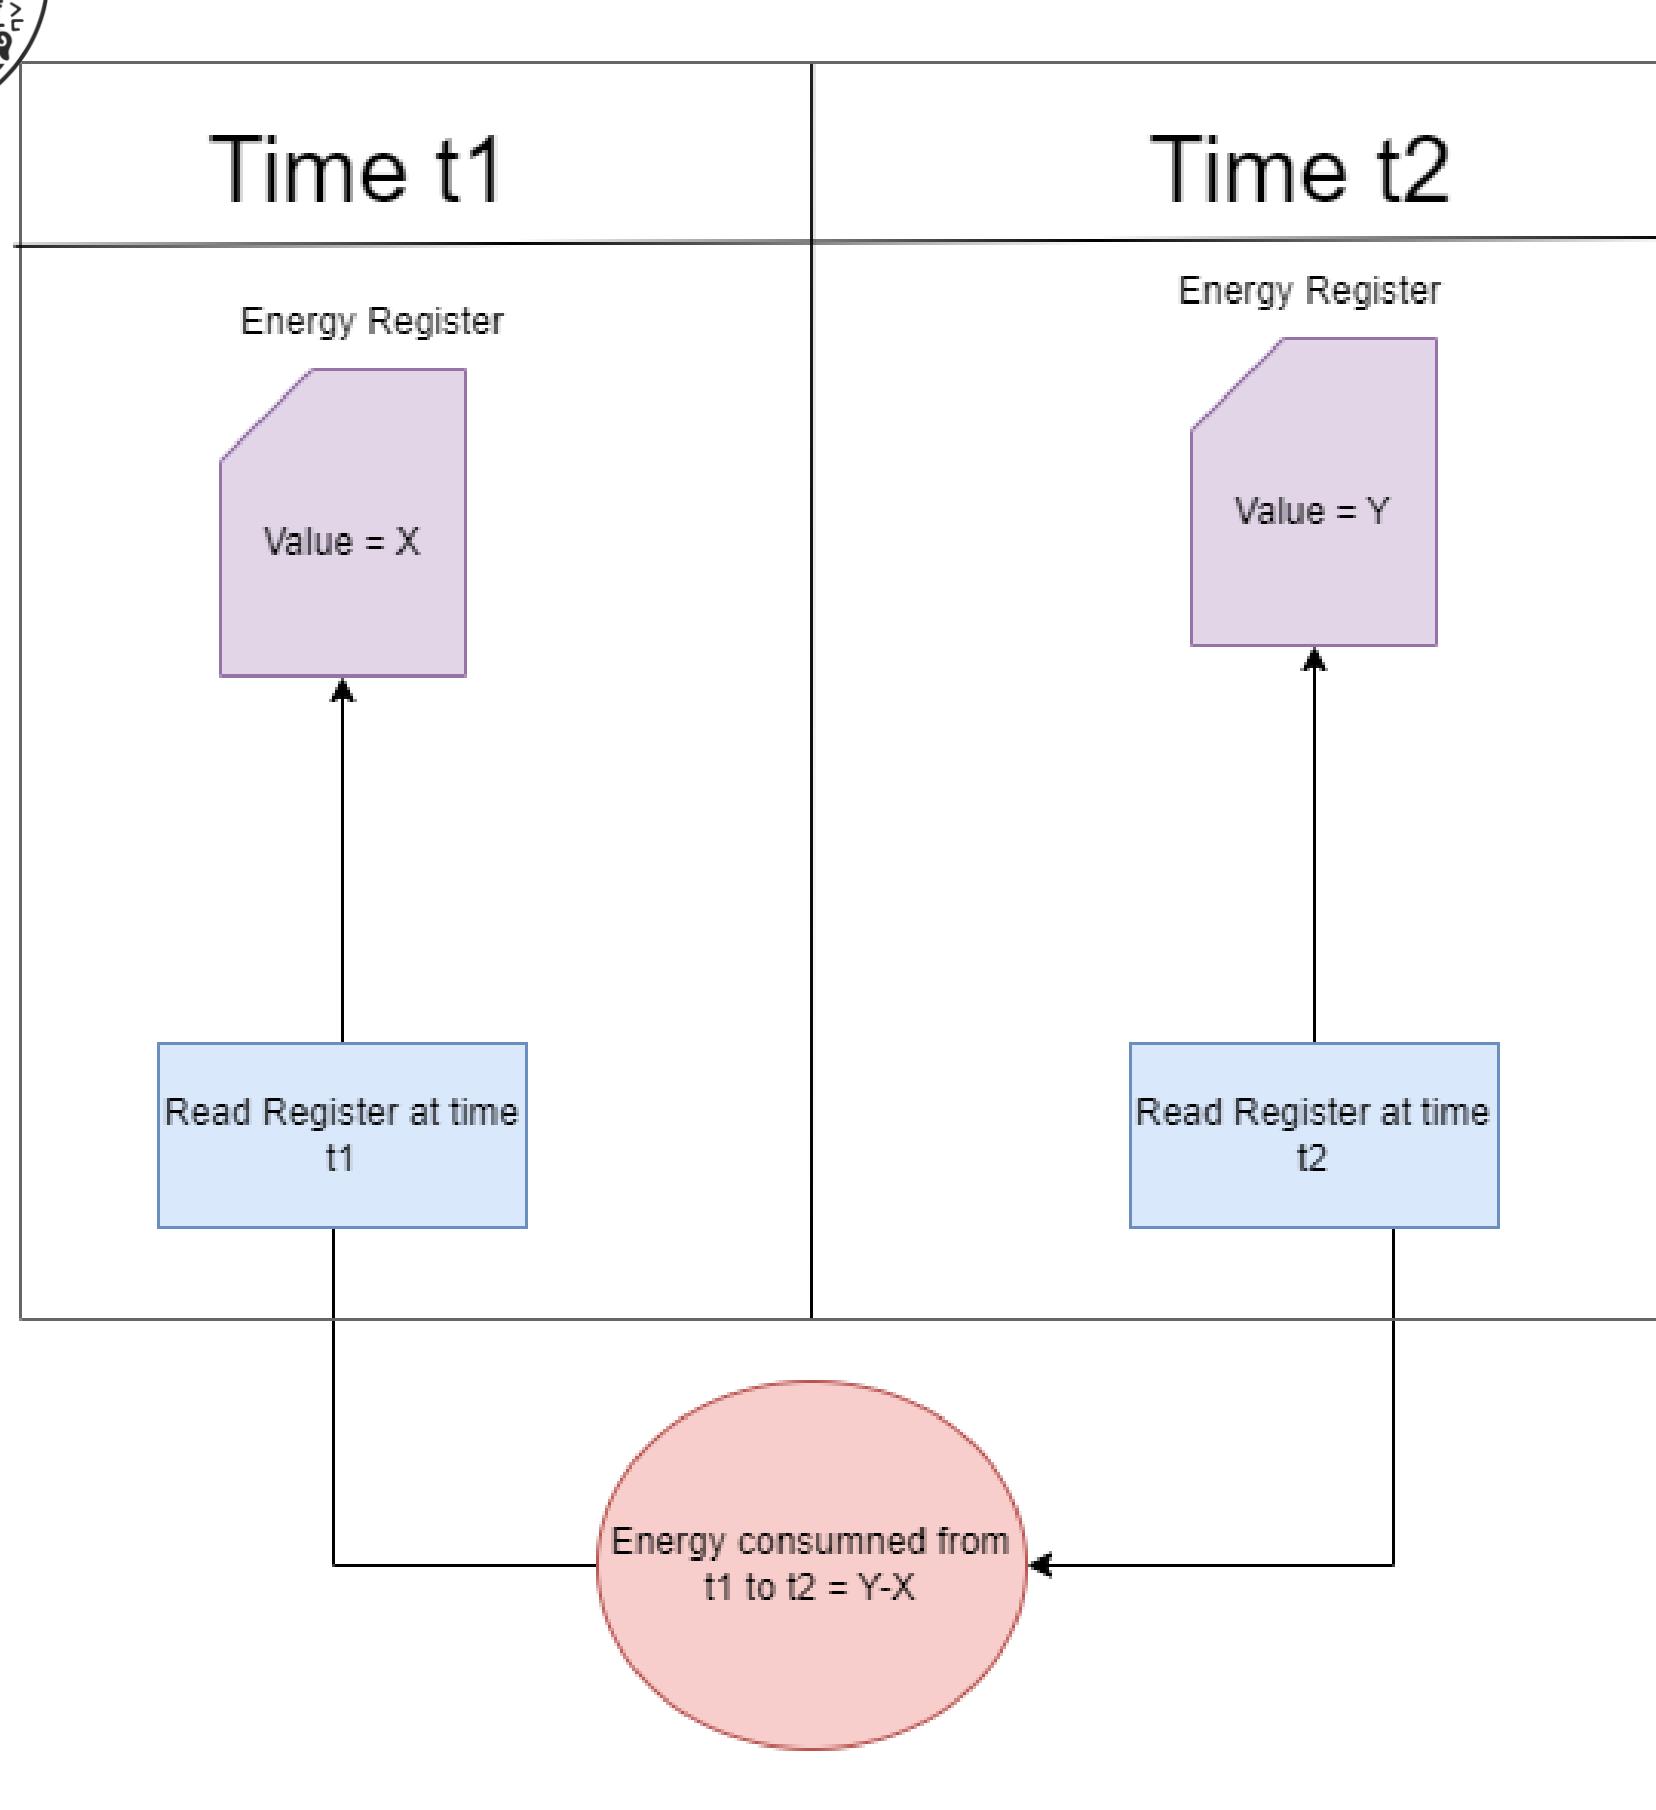
Pipeline





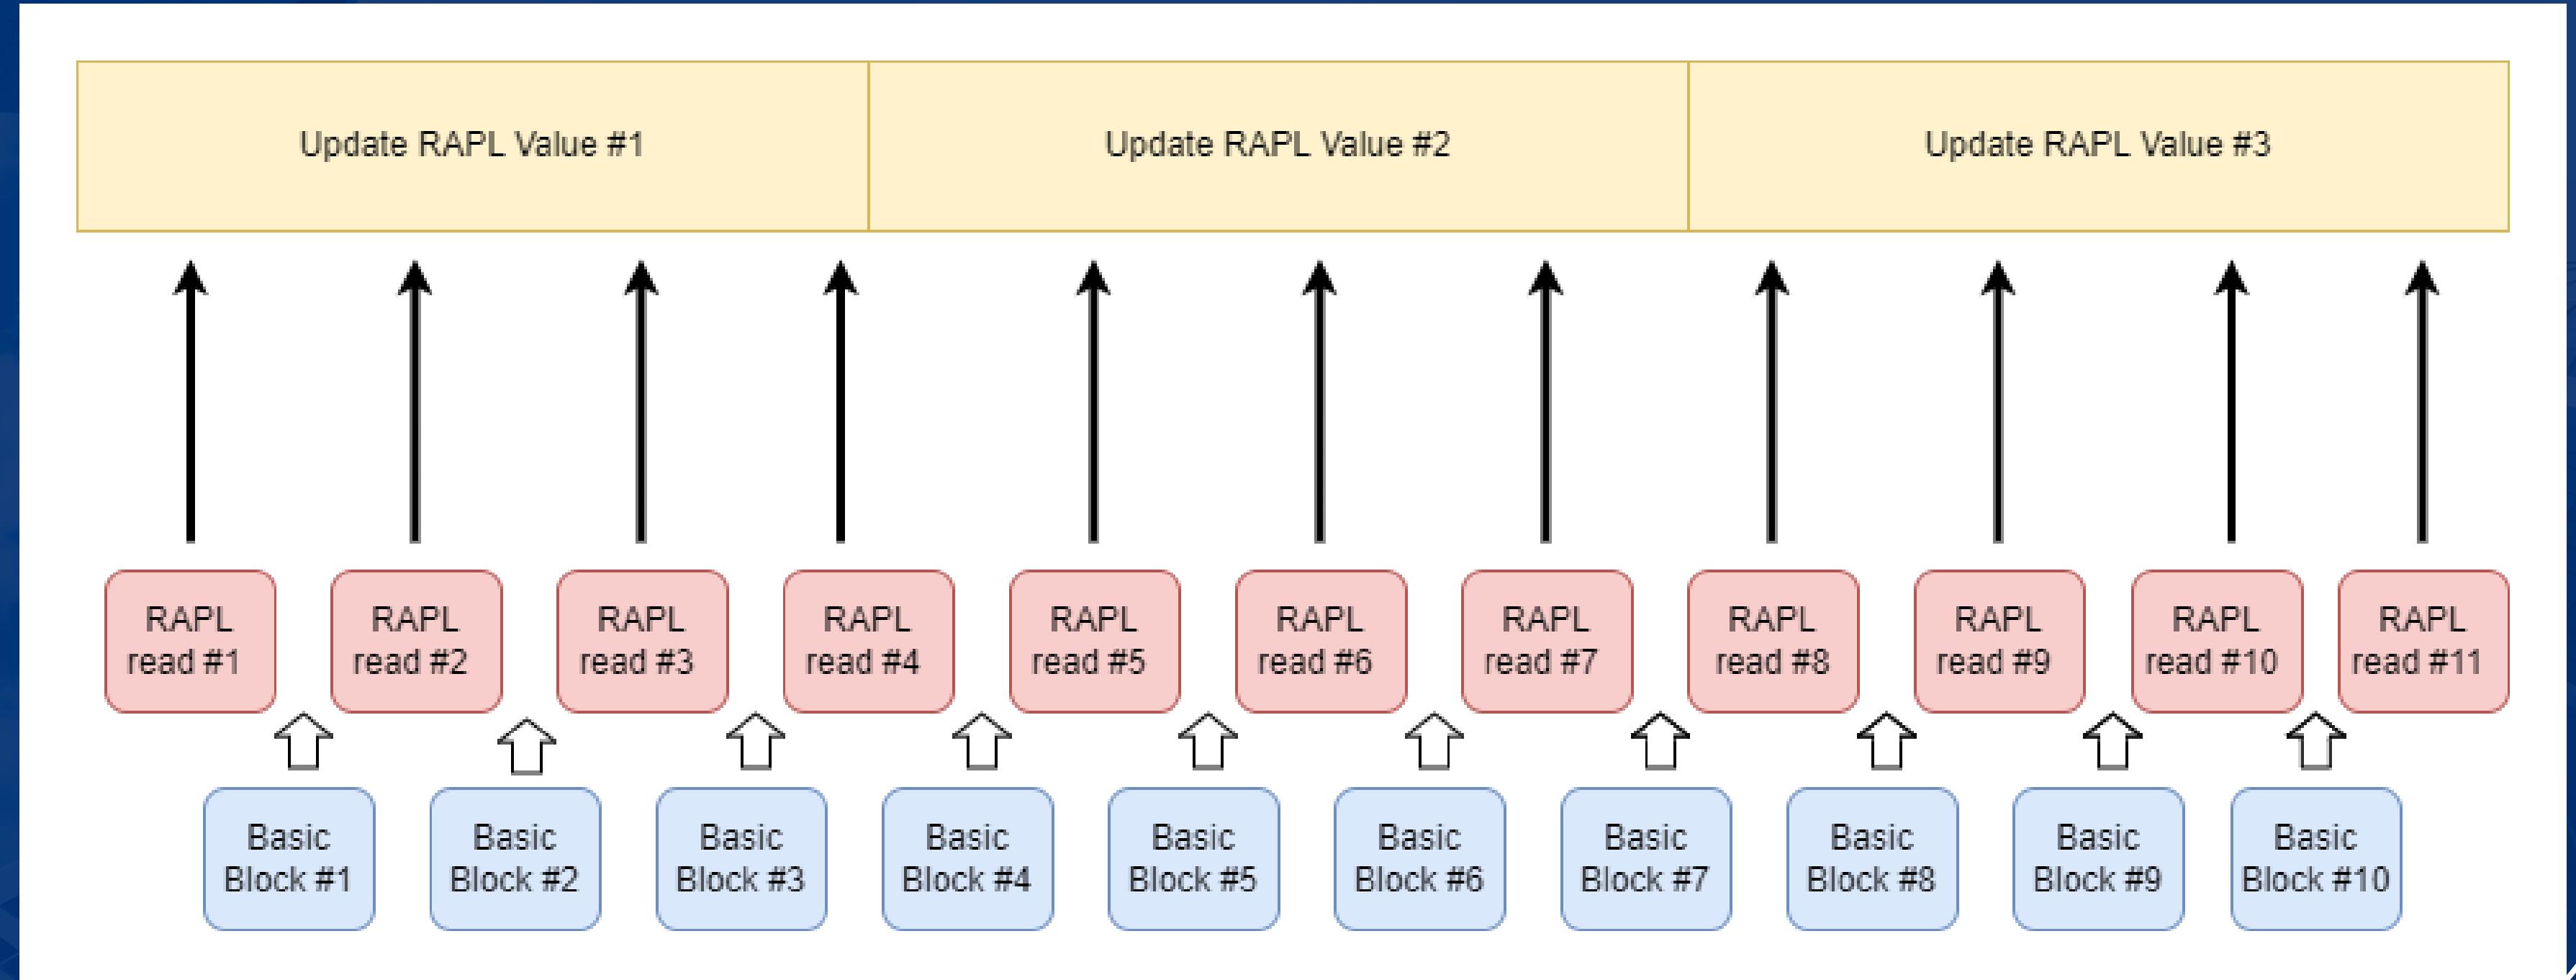
First approach : individual basic block measurement using RAPL calls

- Create an unique ID for each BB
- Add RAPL call function before every BB
- Measure energies through subtraction
- Correspond energy to BB using ID





Problem 1: Ineffectiveness of RAPL read granularity





Solution 1: Split energy based on sum of execution time of assembly commands of each BB

W_{iz} = execution time of command z of BB i

t_i = execution time of BBi (in CPU cycles)

X = energy of multiple BBs

Final energy of BB i=

$$t_i = \sum_{z=1}^M W_{iz}$$

[1]

$$\sum_{i=1}^N t_i = T$$

[2]

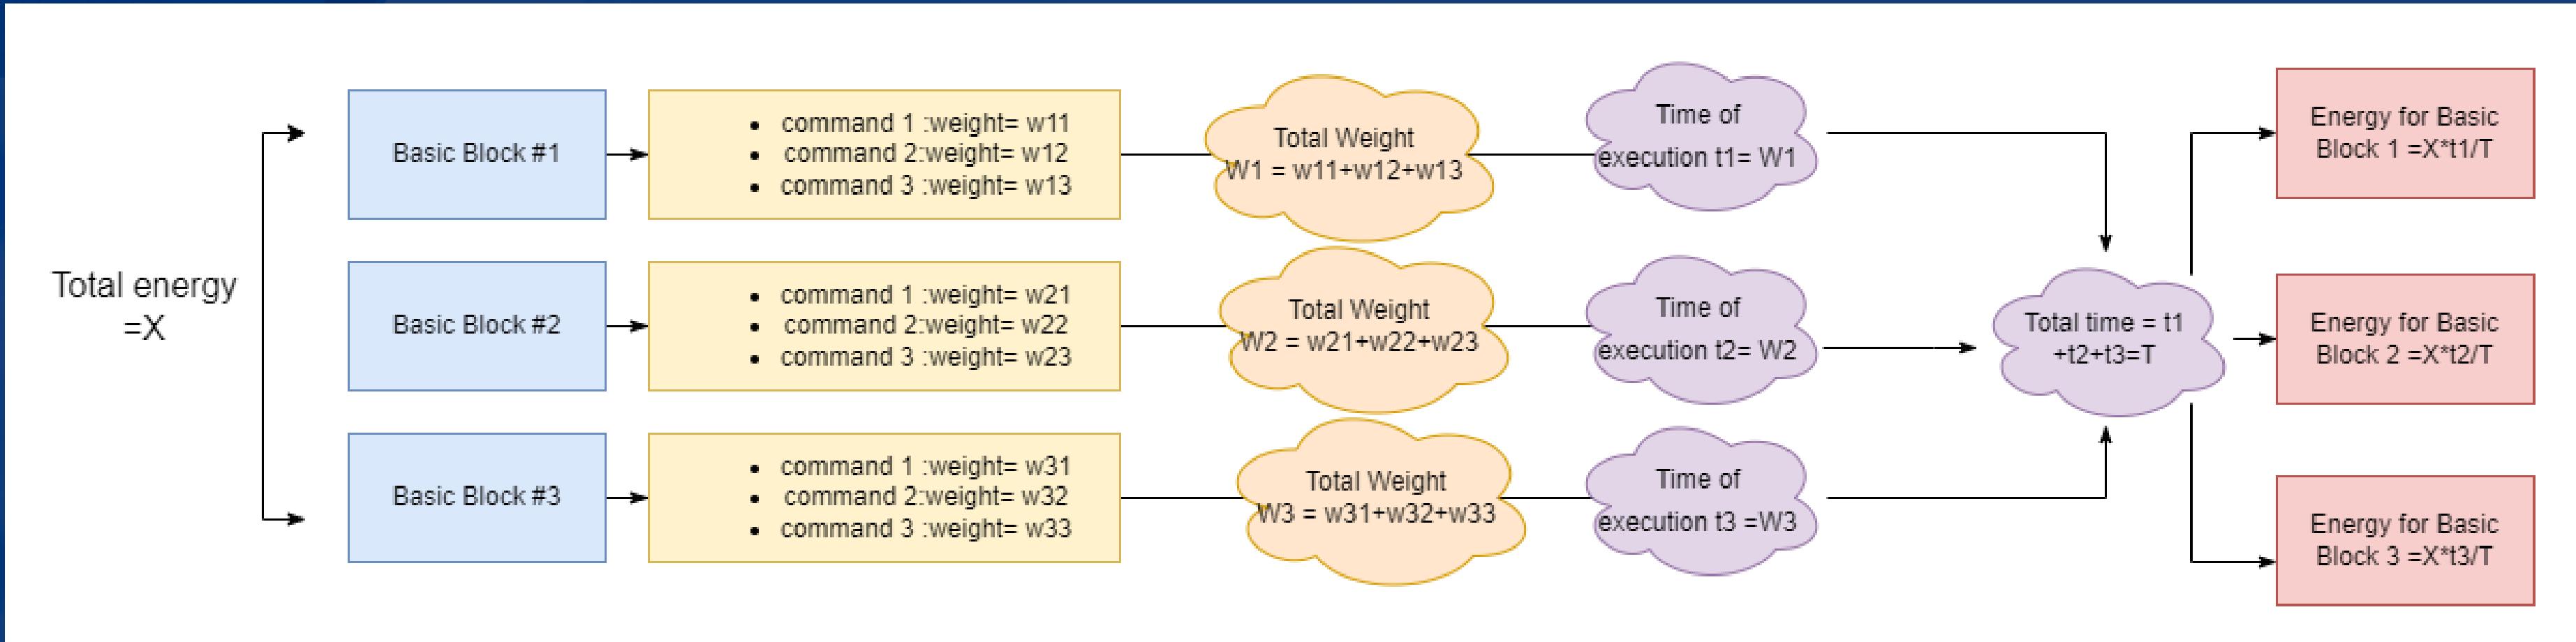
$$\frac{X * t_i}{T}$$

[1] Tiwari, S. Malik, A. Wolfe και M.T. C. Lee.
Instruction level power analysis and optimization of software. Proceedings of 9th International Conference on VLSI Design, σελίδες 326–328, 1996

[2] Lev Mukhanov, Dimitrios S. Nikolopoulos και Bronis R. De Supinski. ALEA: Fine-Grain Energy Profiling with Basic Block Sampling. 2015 International Conference on Parallel Architecture and Compilation (PACT), σελίδες 87–98, 2015



Solution 1: Split energy based on sum of execution time of assembly commands of each BB





Problem 2: Cost of RAPL overshadowing cost of Basic Block

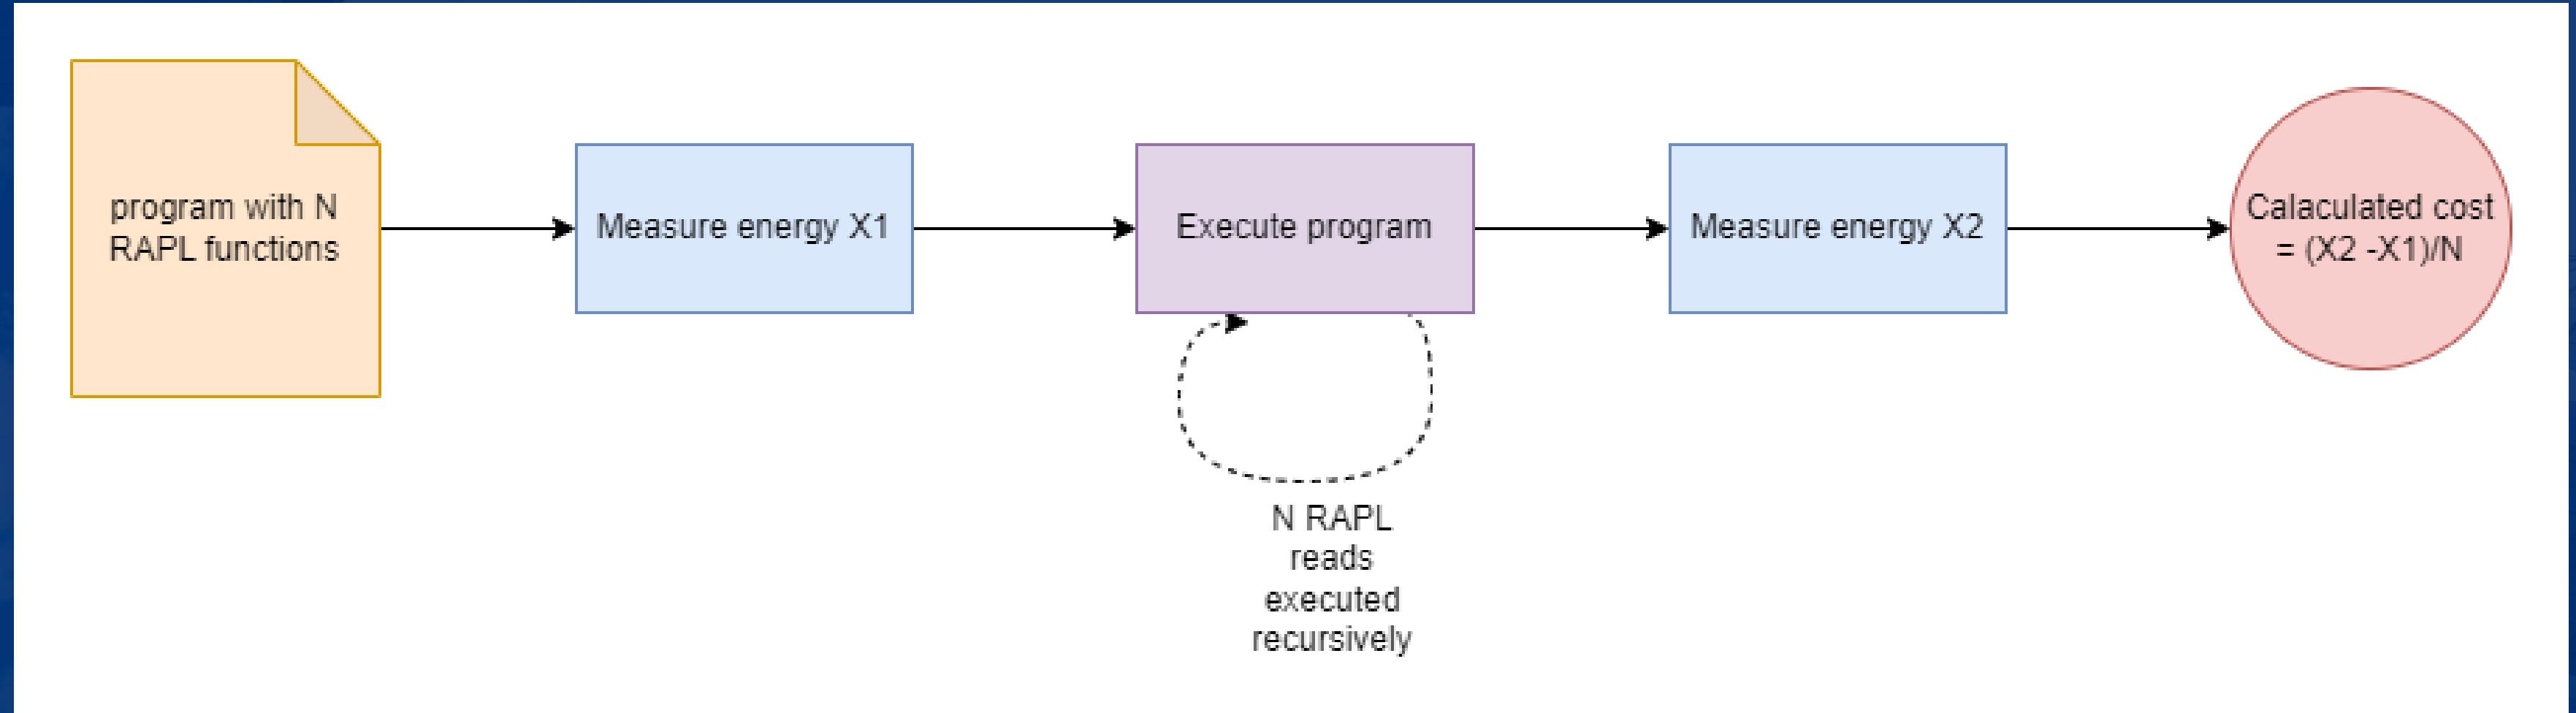
RAPL read function introduces energy overhead of up to 20X times the energy cost of BB

but

It is always the same , so we can calculate it and subtract it

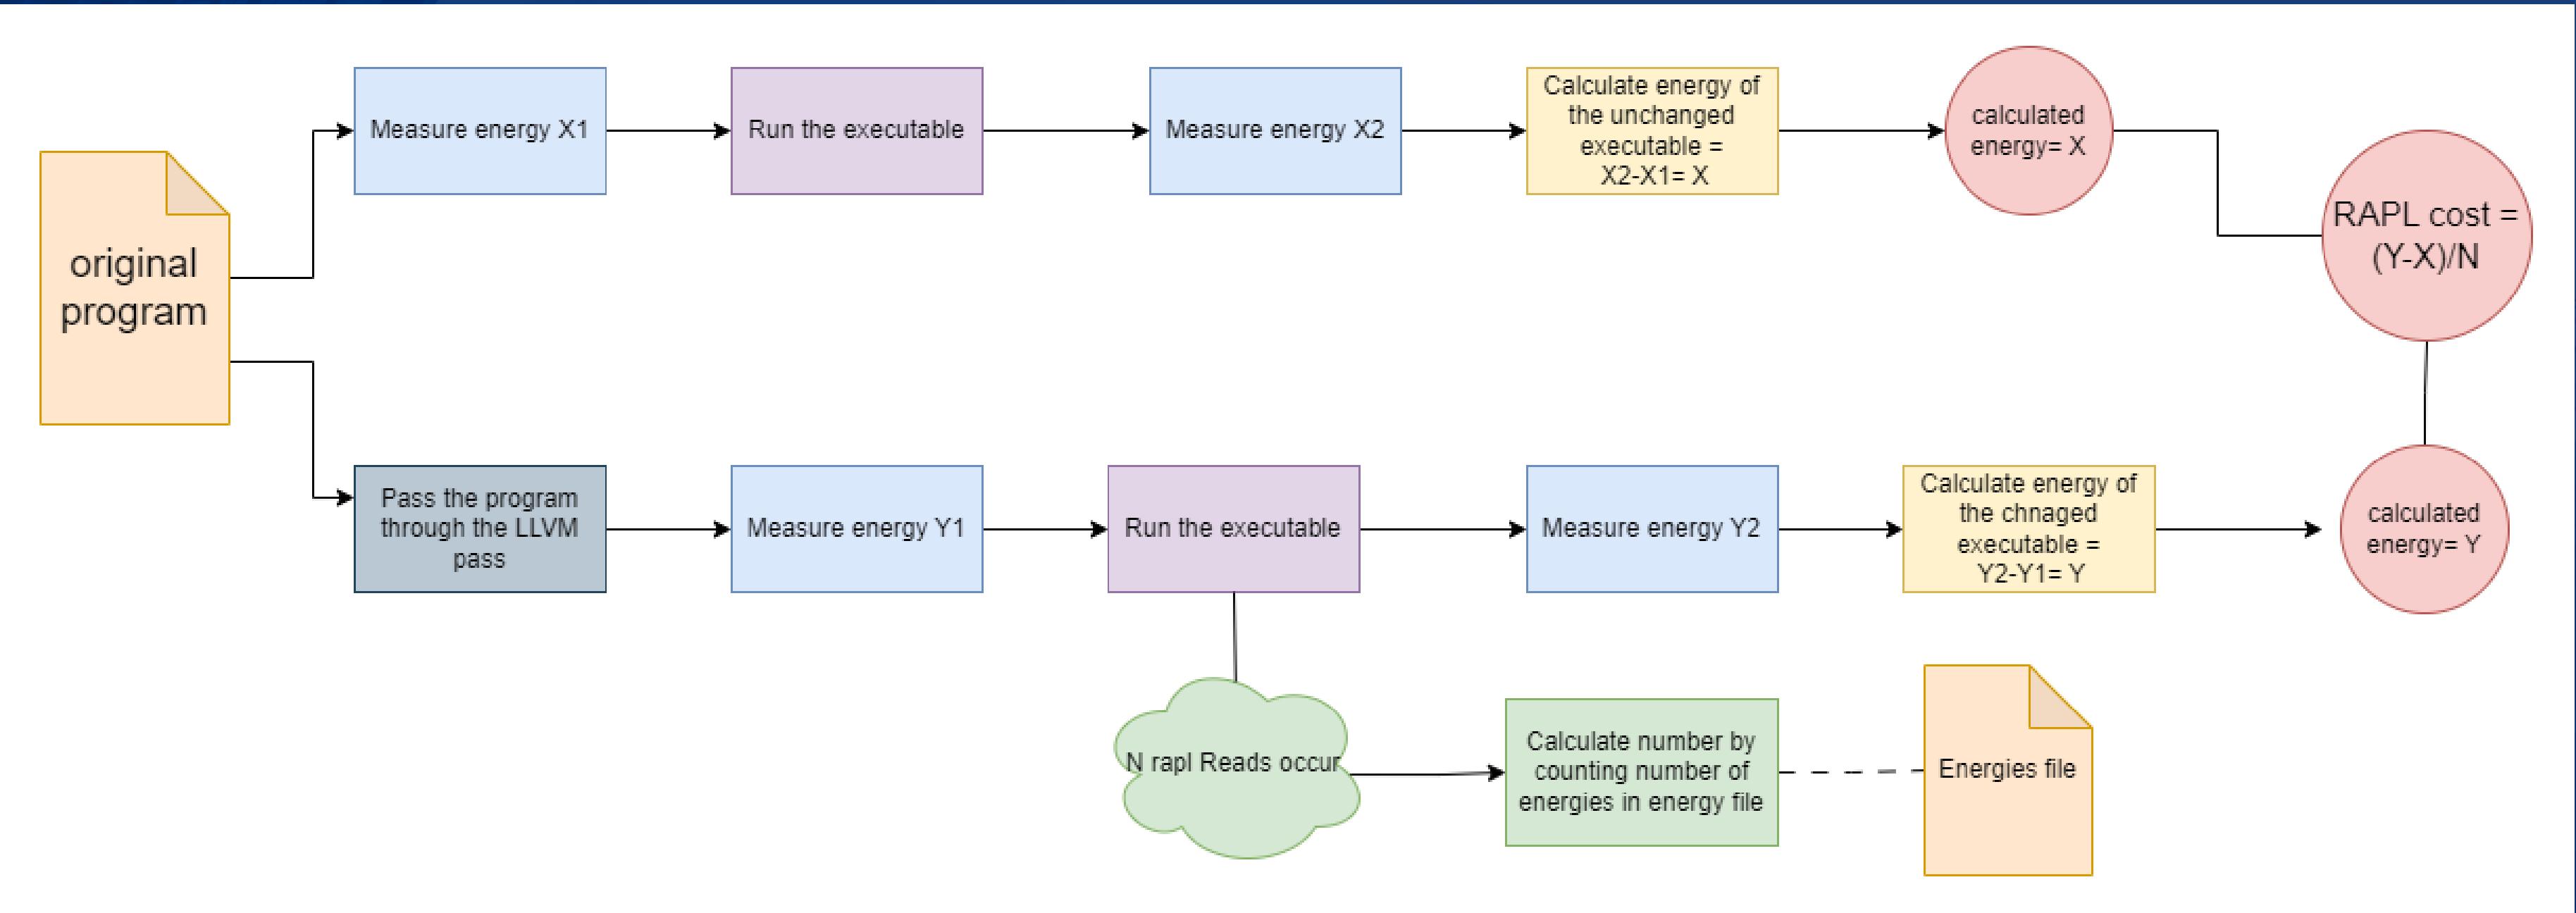


Solution 2: Subtract fixed cost of RAPL read





Solution 2: Subtract fixed cost of RAPL read



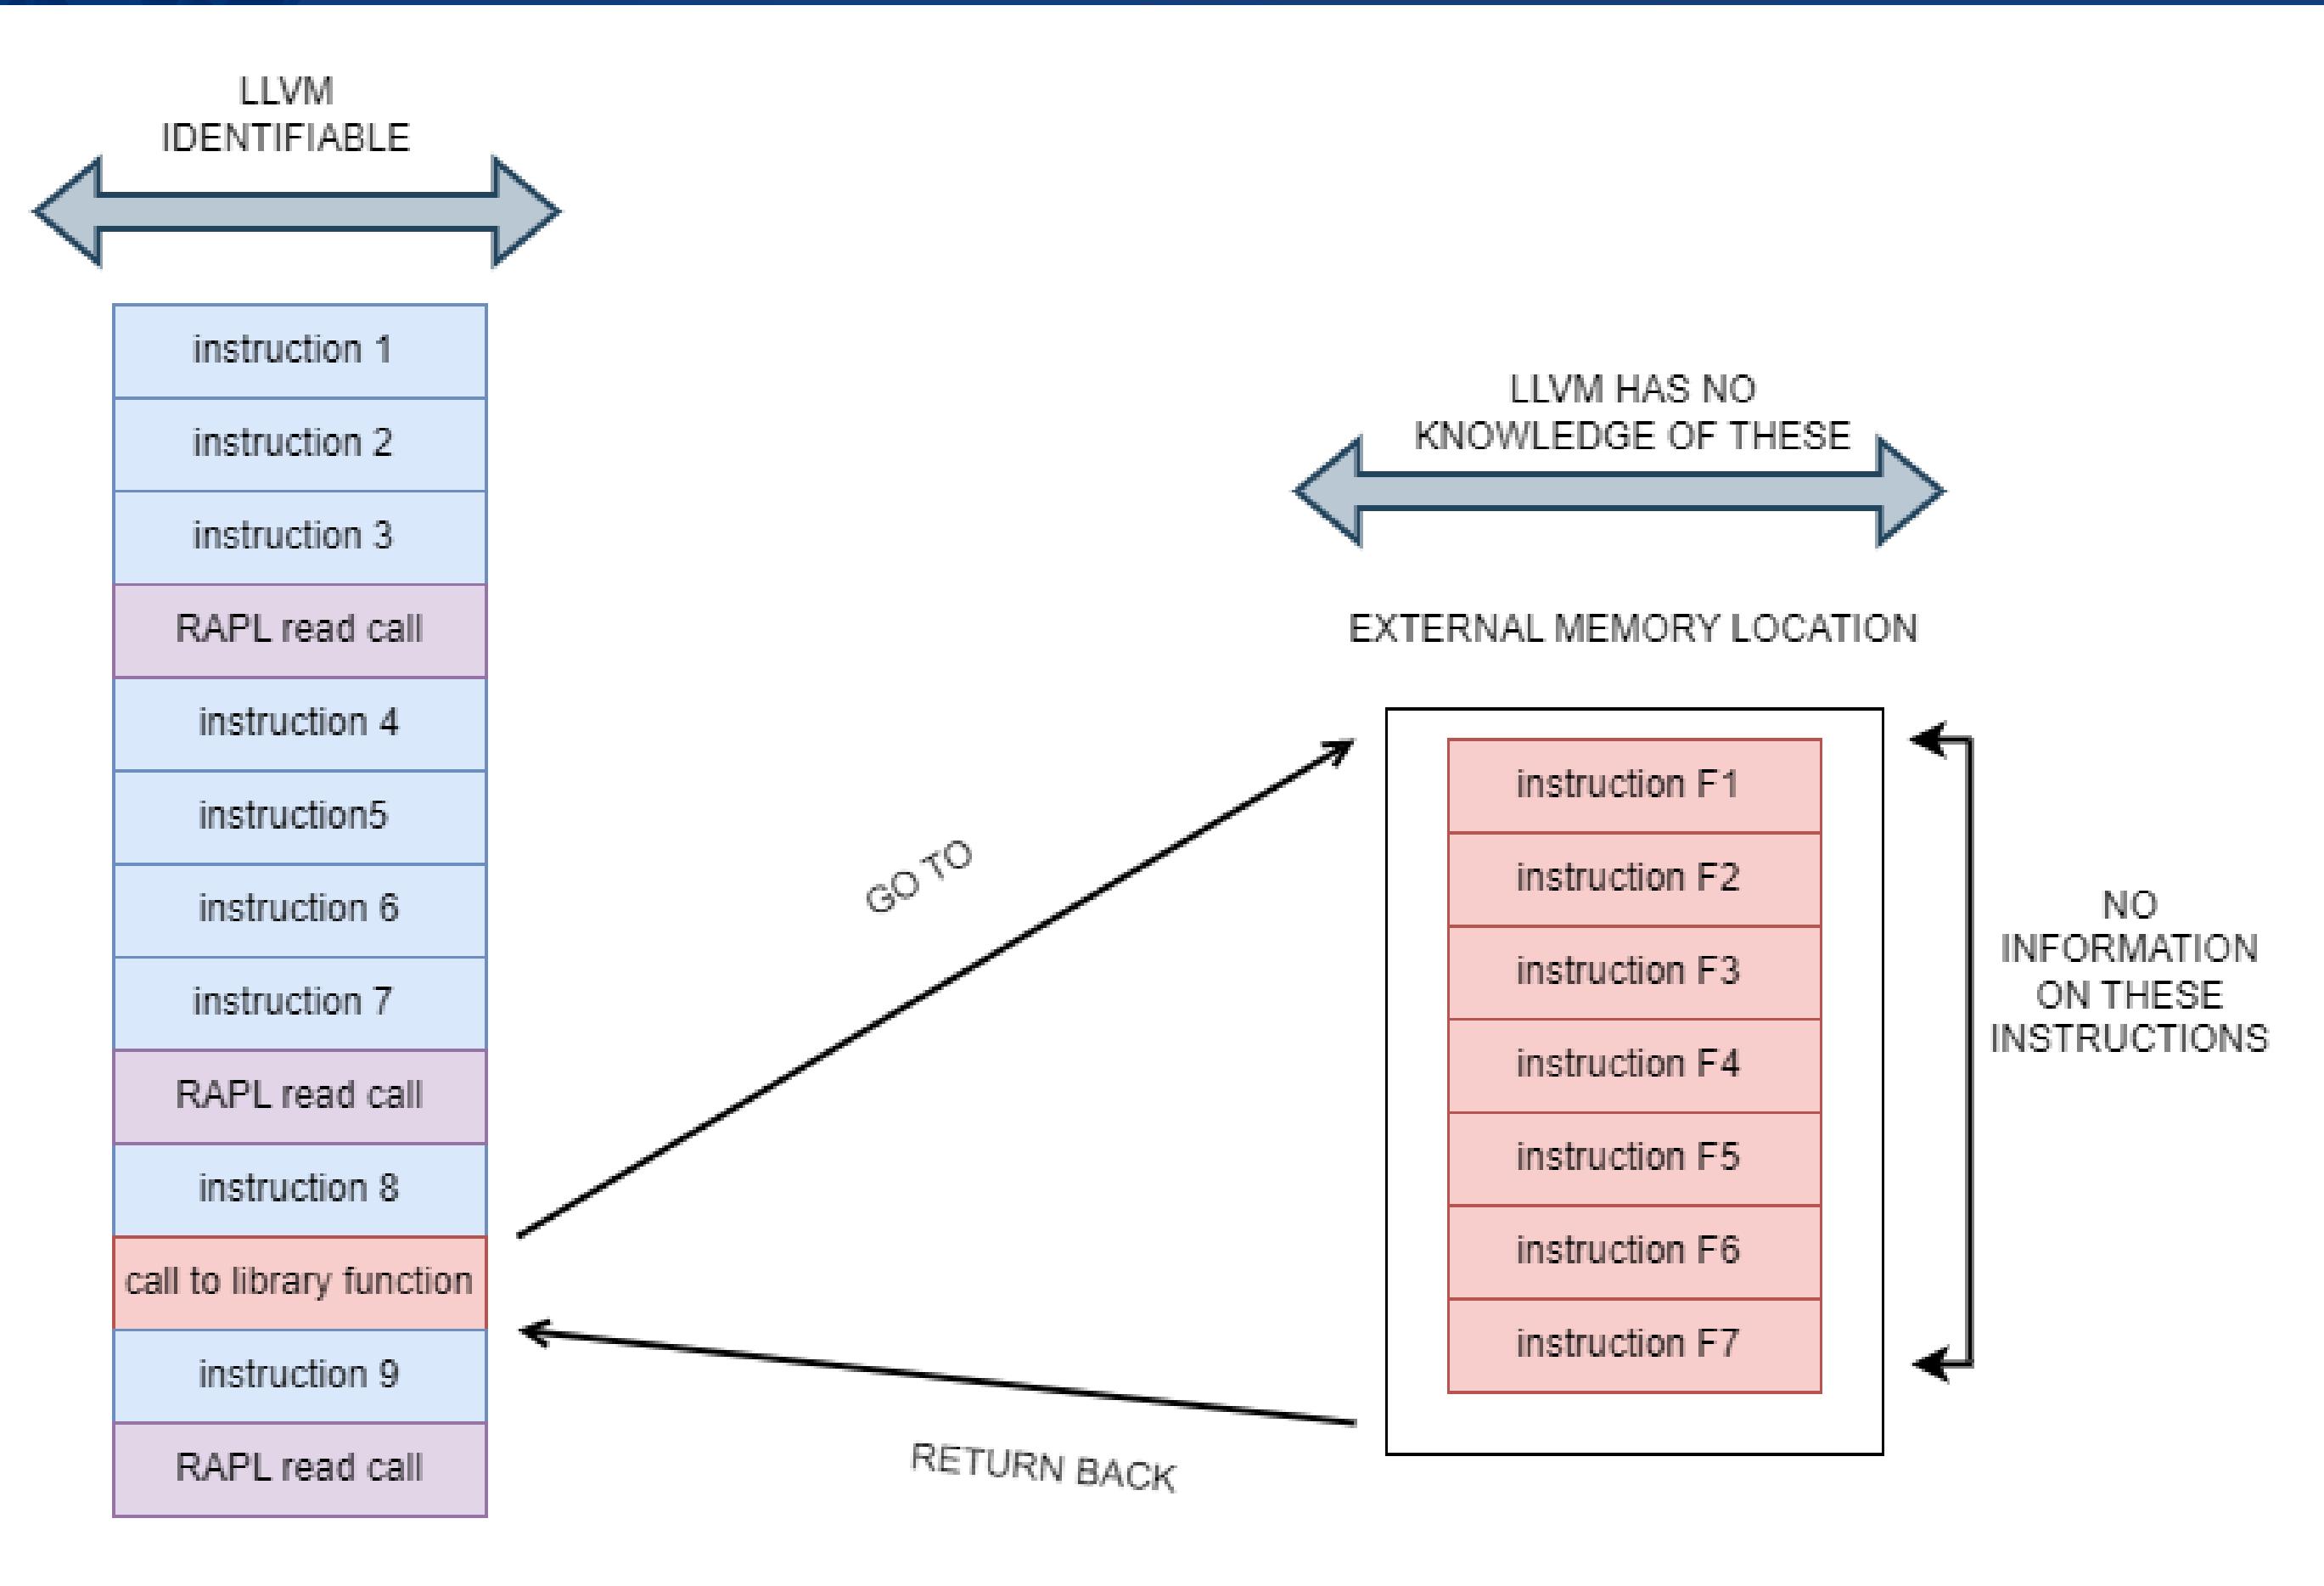


Problem 3: Lost library functions

Library functions are added during linking time and contain bulk of the code

so

LLVM pass cannot analyze their BBs or add RAPL calls





Solution 3: find the code of the library functions



Could not find viable solution so we were lead to
Approach 2

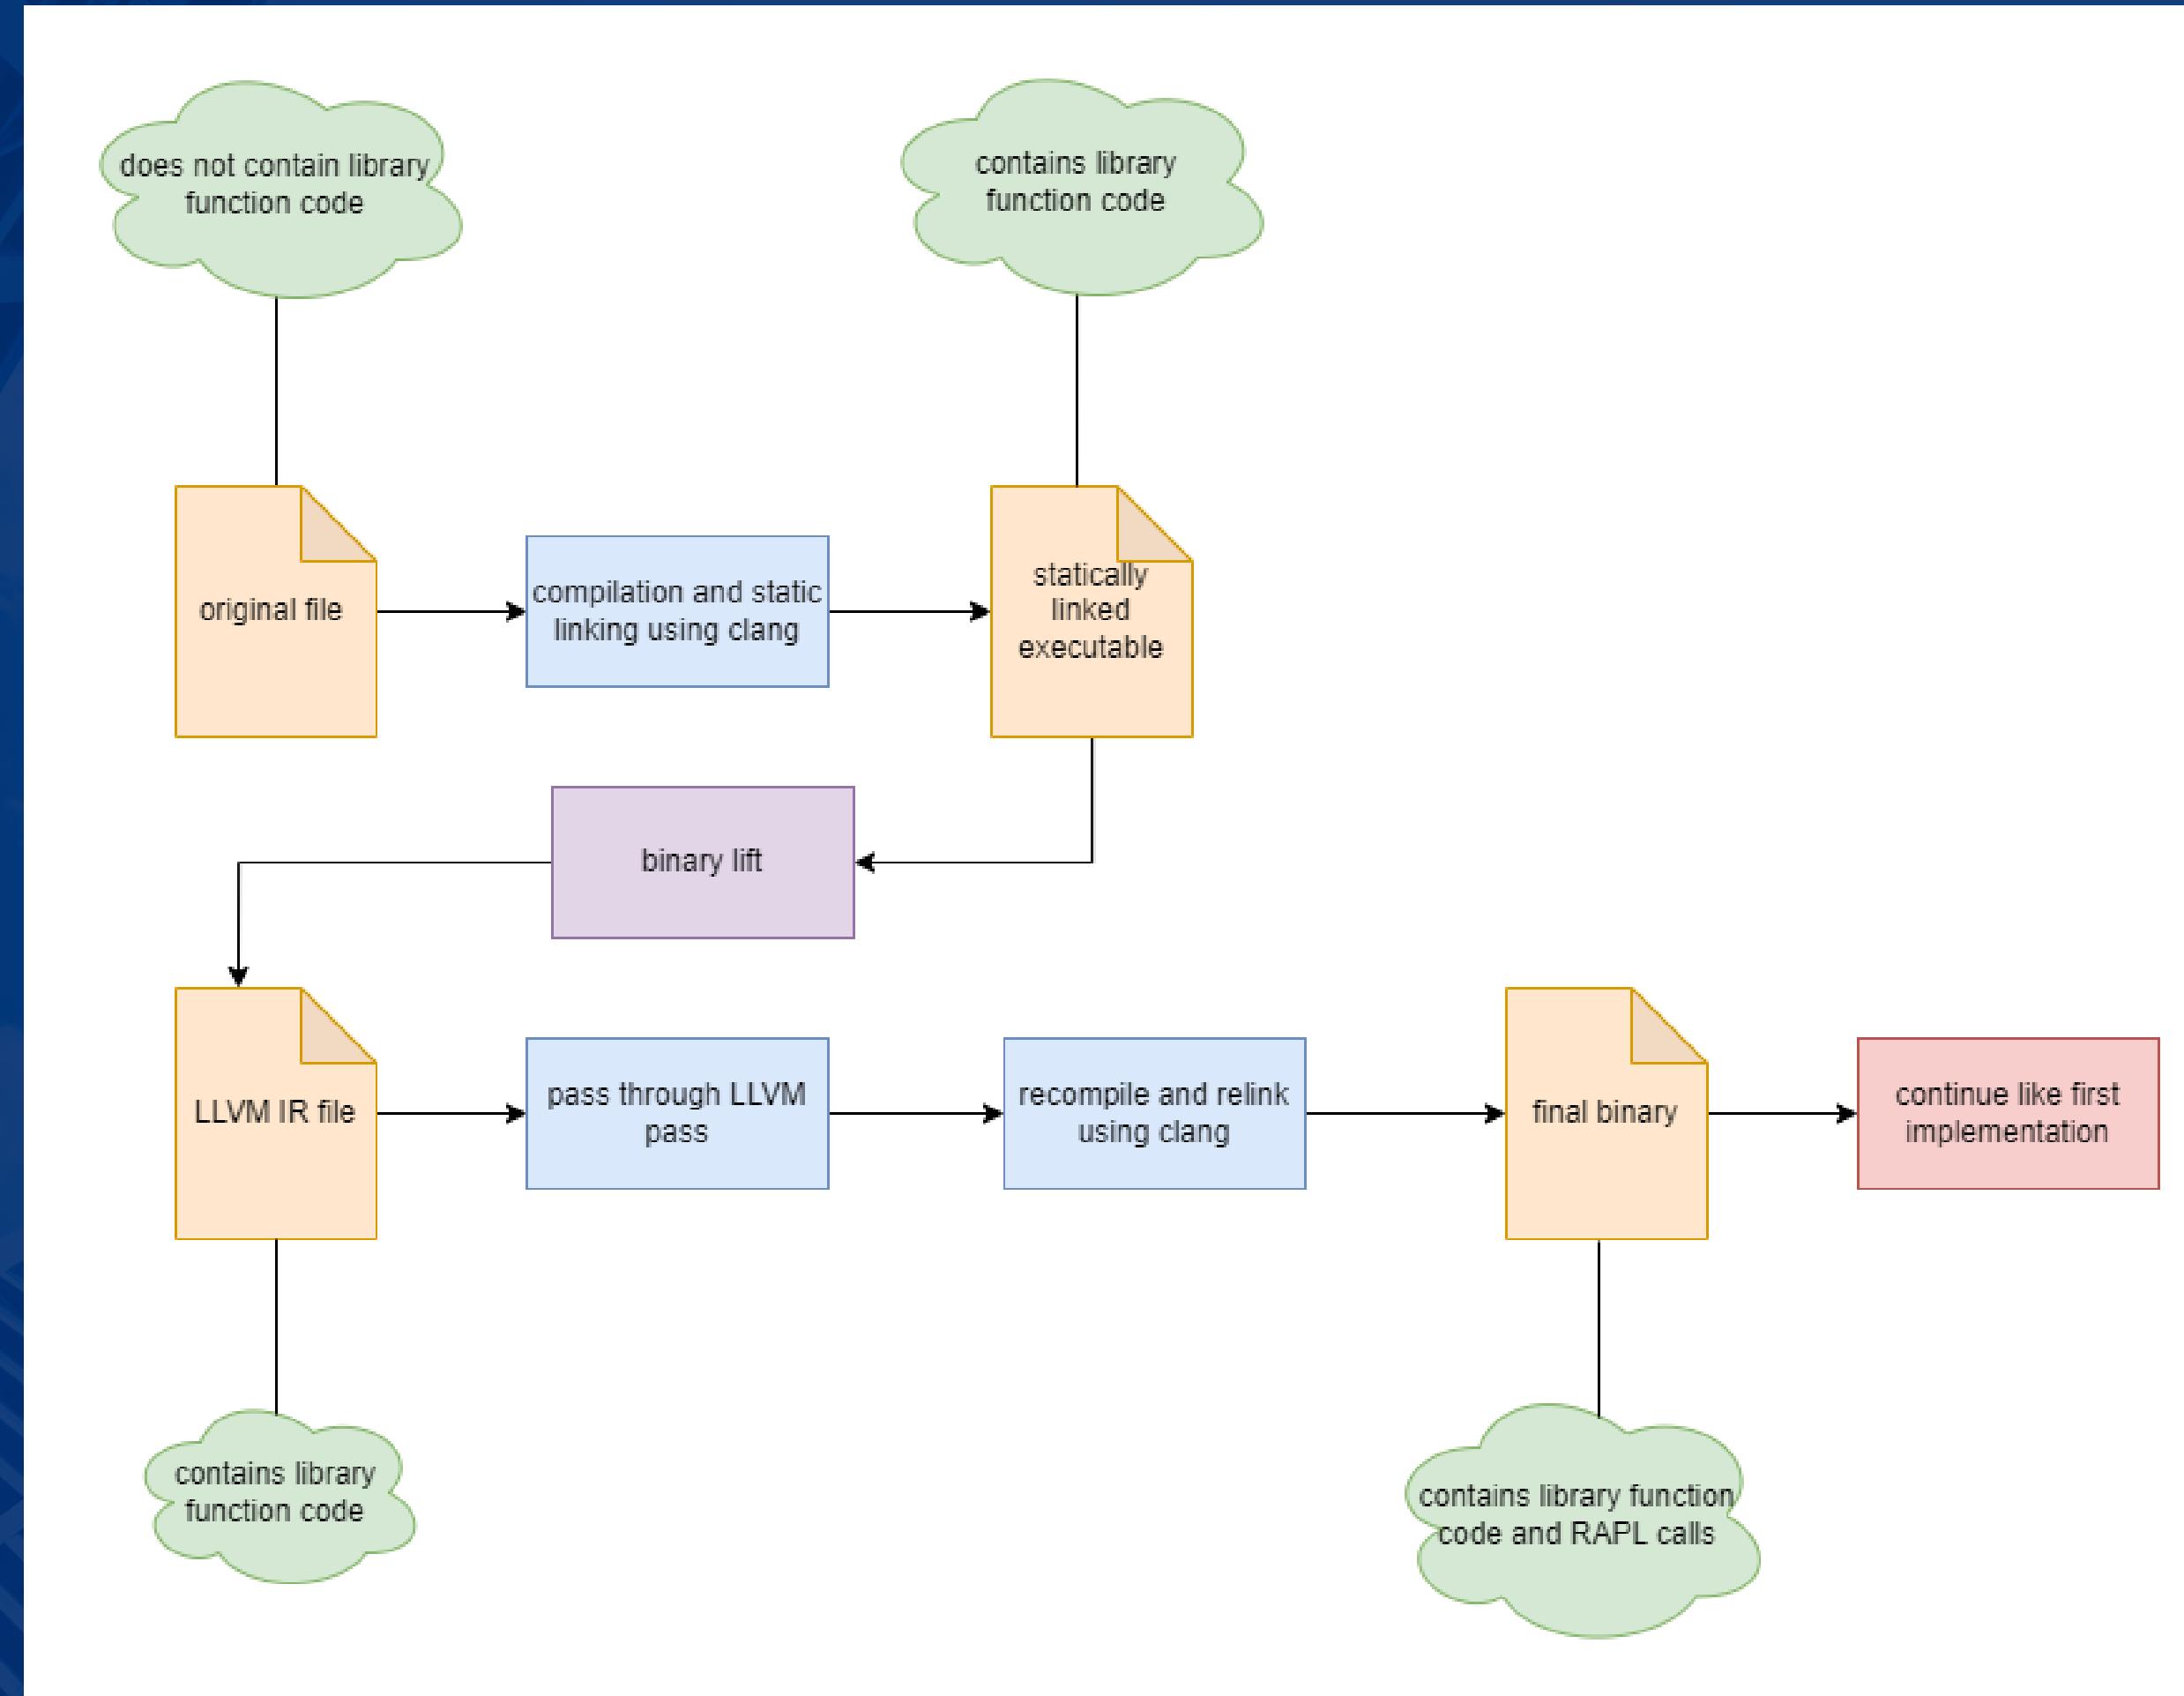


Second approach : binary lift of linked executable to LLVM IR

- LLVM pass can only be performed on LLVM IR
- Library function code is added during link time

so

- Transform linked executable back to IR
- Pass it through the LLVM pass





Problem : lifted binary is not executable
after it passes through the LLVM pass

Used 3 different binary lifting tools

Managed to lift executable and pass it through
LLVM pass

Resulting executable always failed with
segmentation fault

**Must find other way to have access to
library function code ...**



Third approach : Execution tracing using Intel perf

- Insert 3 types of RAPL calls
- Run executable and collect trace
- Remove RAPL call energy and code
- Analyze trace corresponding energy to BBS or functions
- Split functions into BBS
- Split function energy to its BBS
- aggregate duplicate BBS



Problem 1: Remove RAPL code and energy

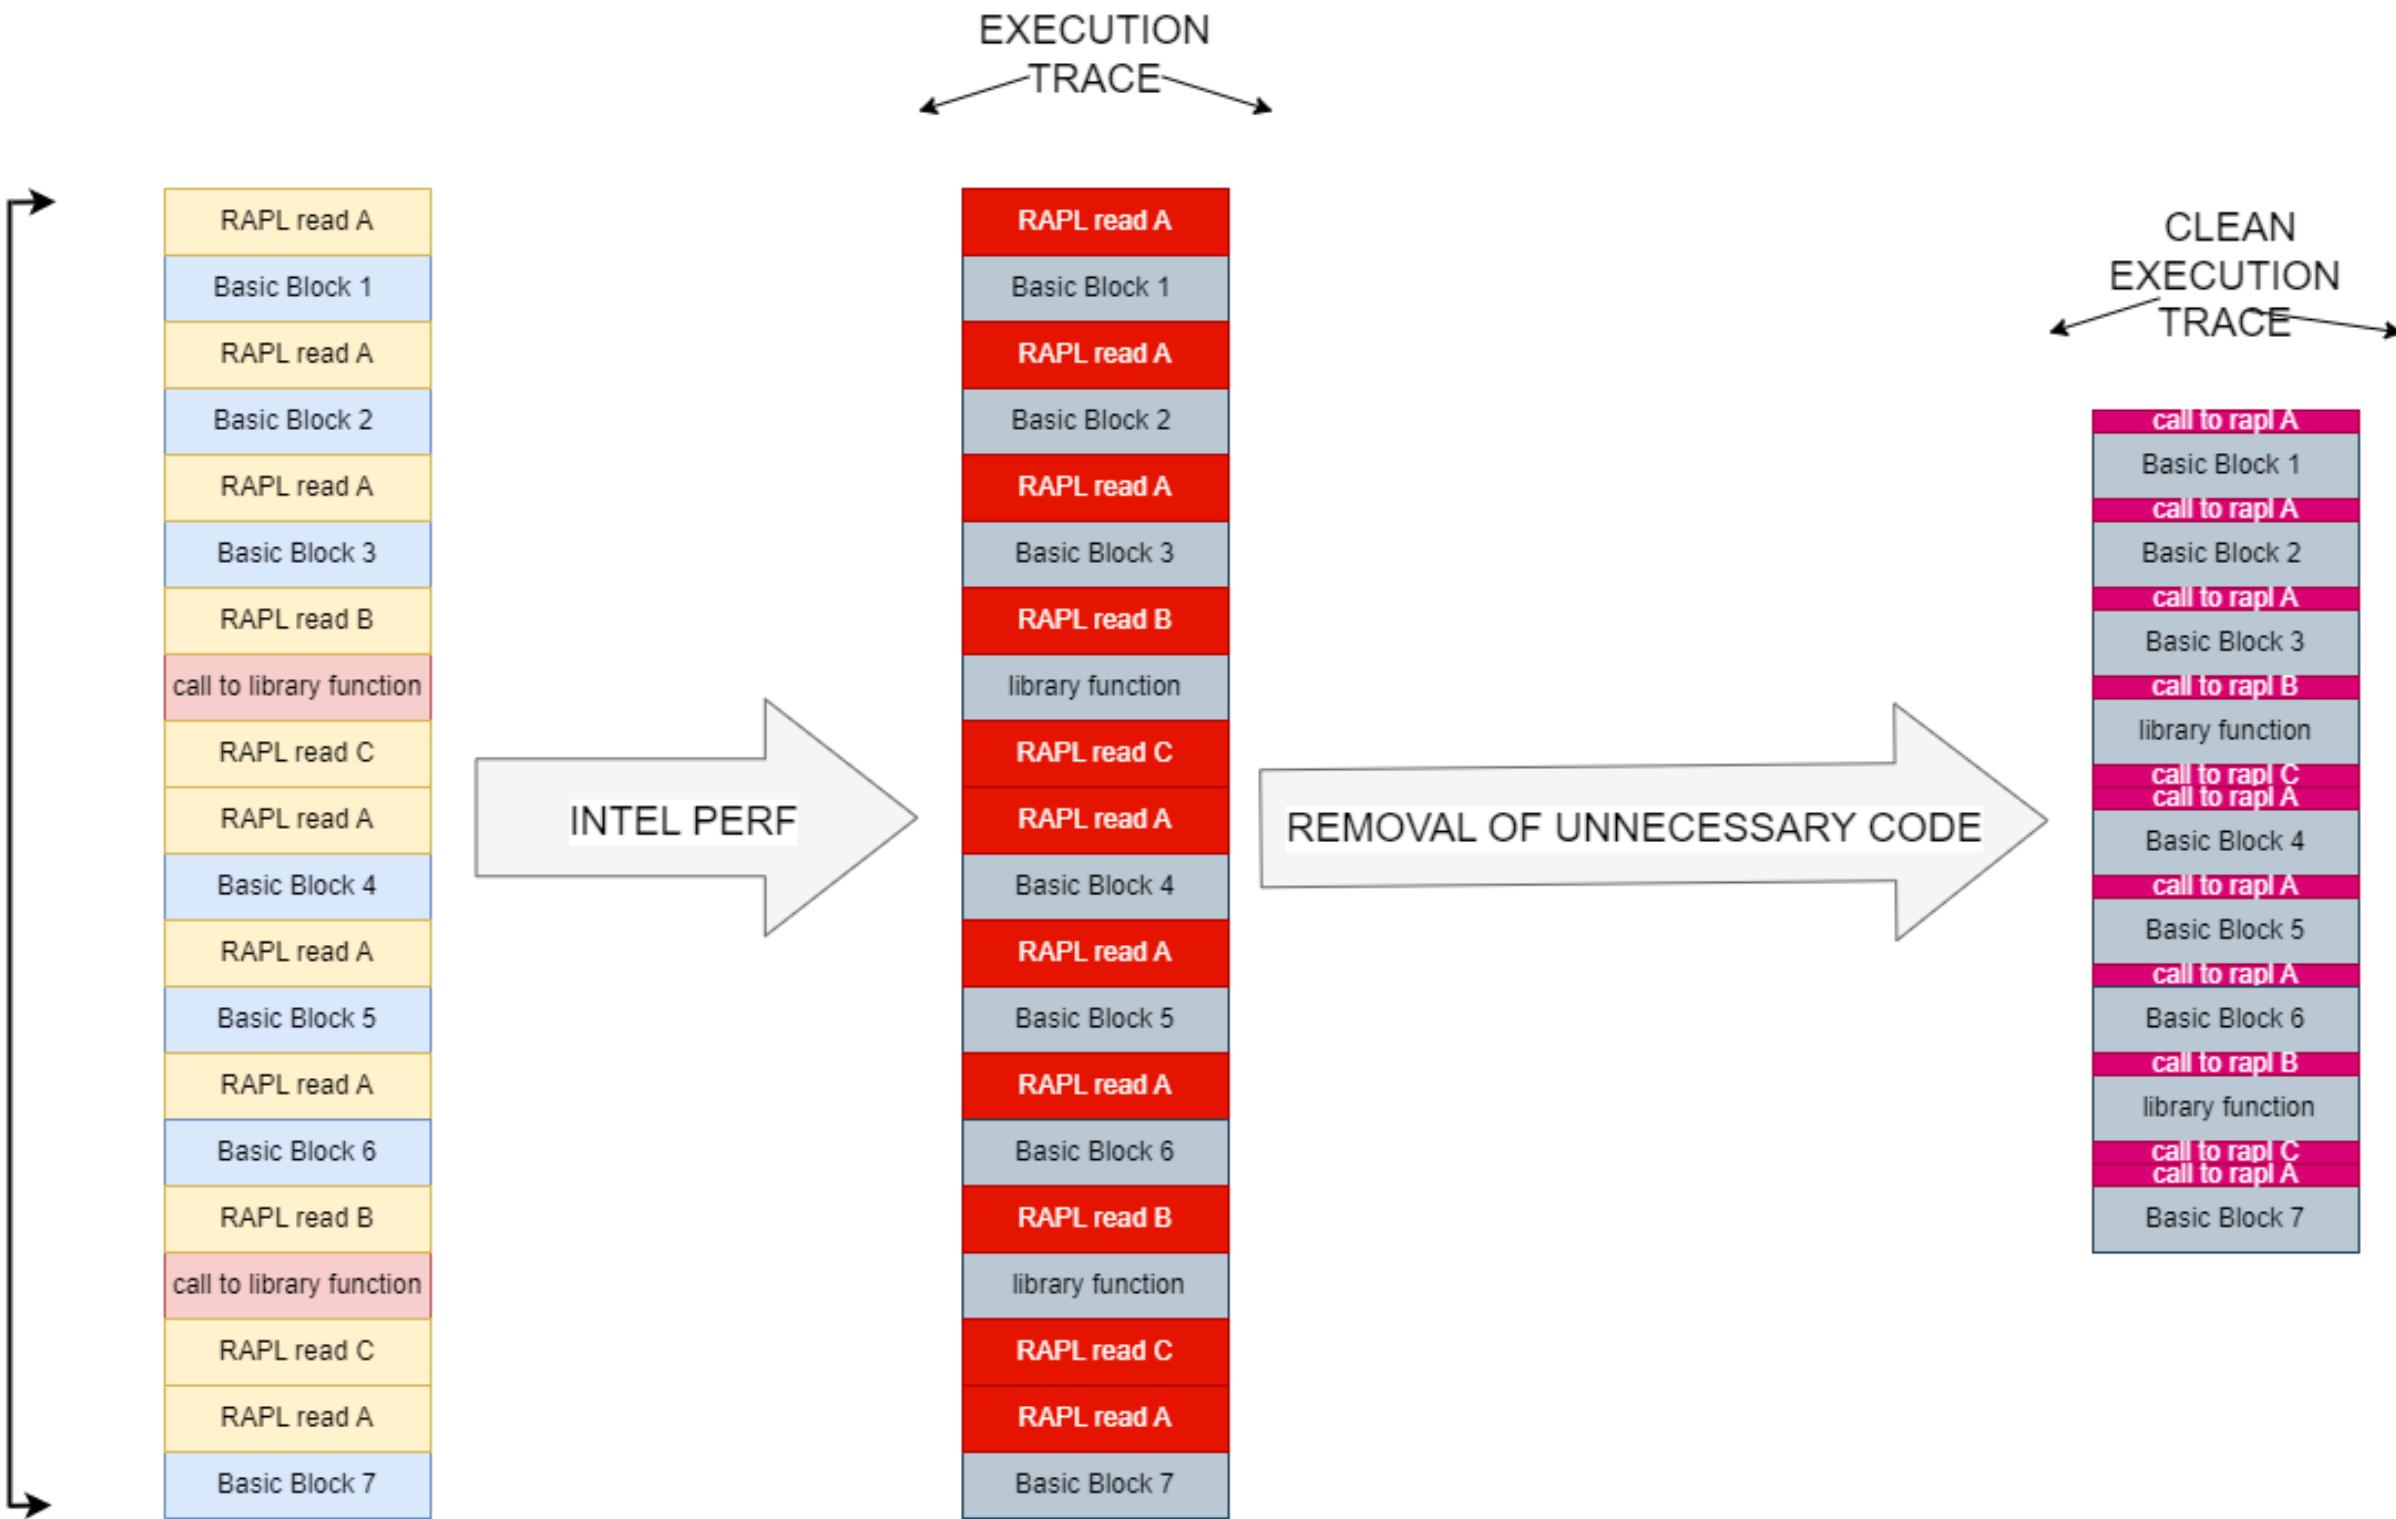
Remove RAPL energy as seen before

and

Remove RAPL code from trace since it is always the same



OUR
PROGRAM
AFTER
GOING
THROUGH
THE LLVM
PASS





Problem 2: Split energy of functions into their BBS

Split functions into BBs using branch or call instructions as break points



Use method based on execution times of instructions to split energy into newly created BBs



Final Dataset

- 3828 unique Basic Blocks from >565K total
- 24 C benchmarks
- Average energy : 0.64 energy units
- Mean error : 2.63%



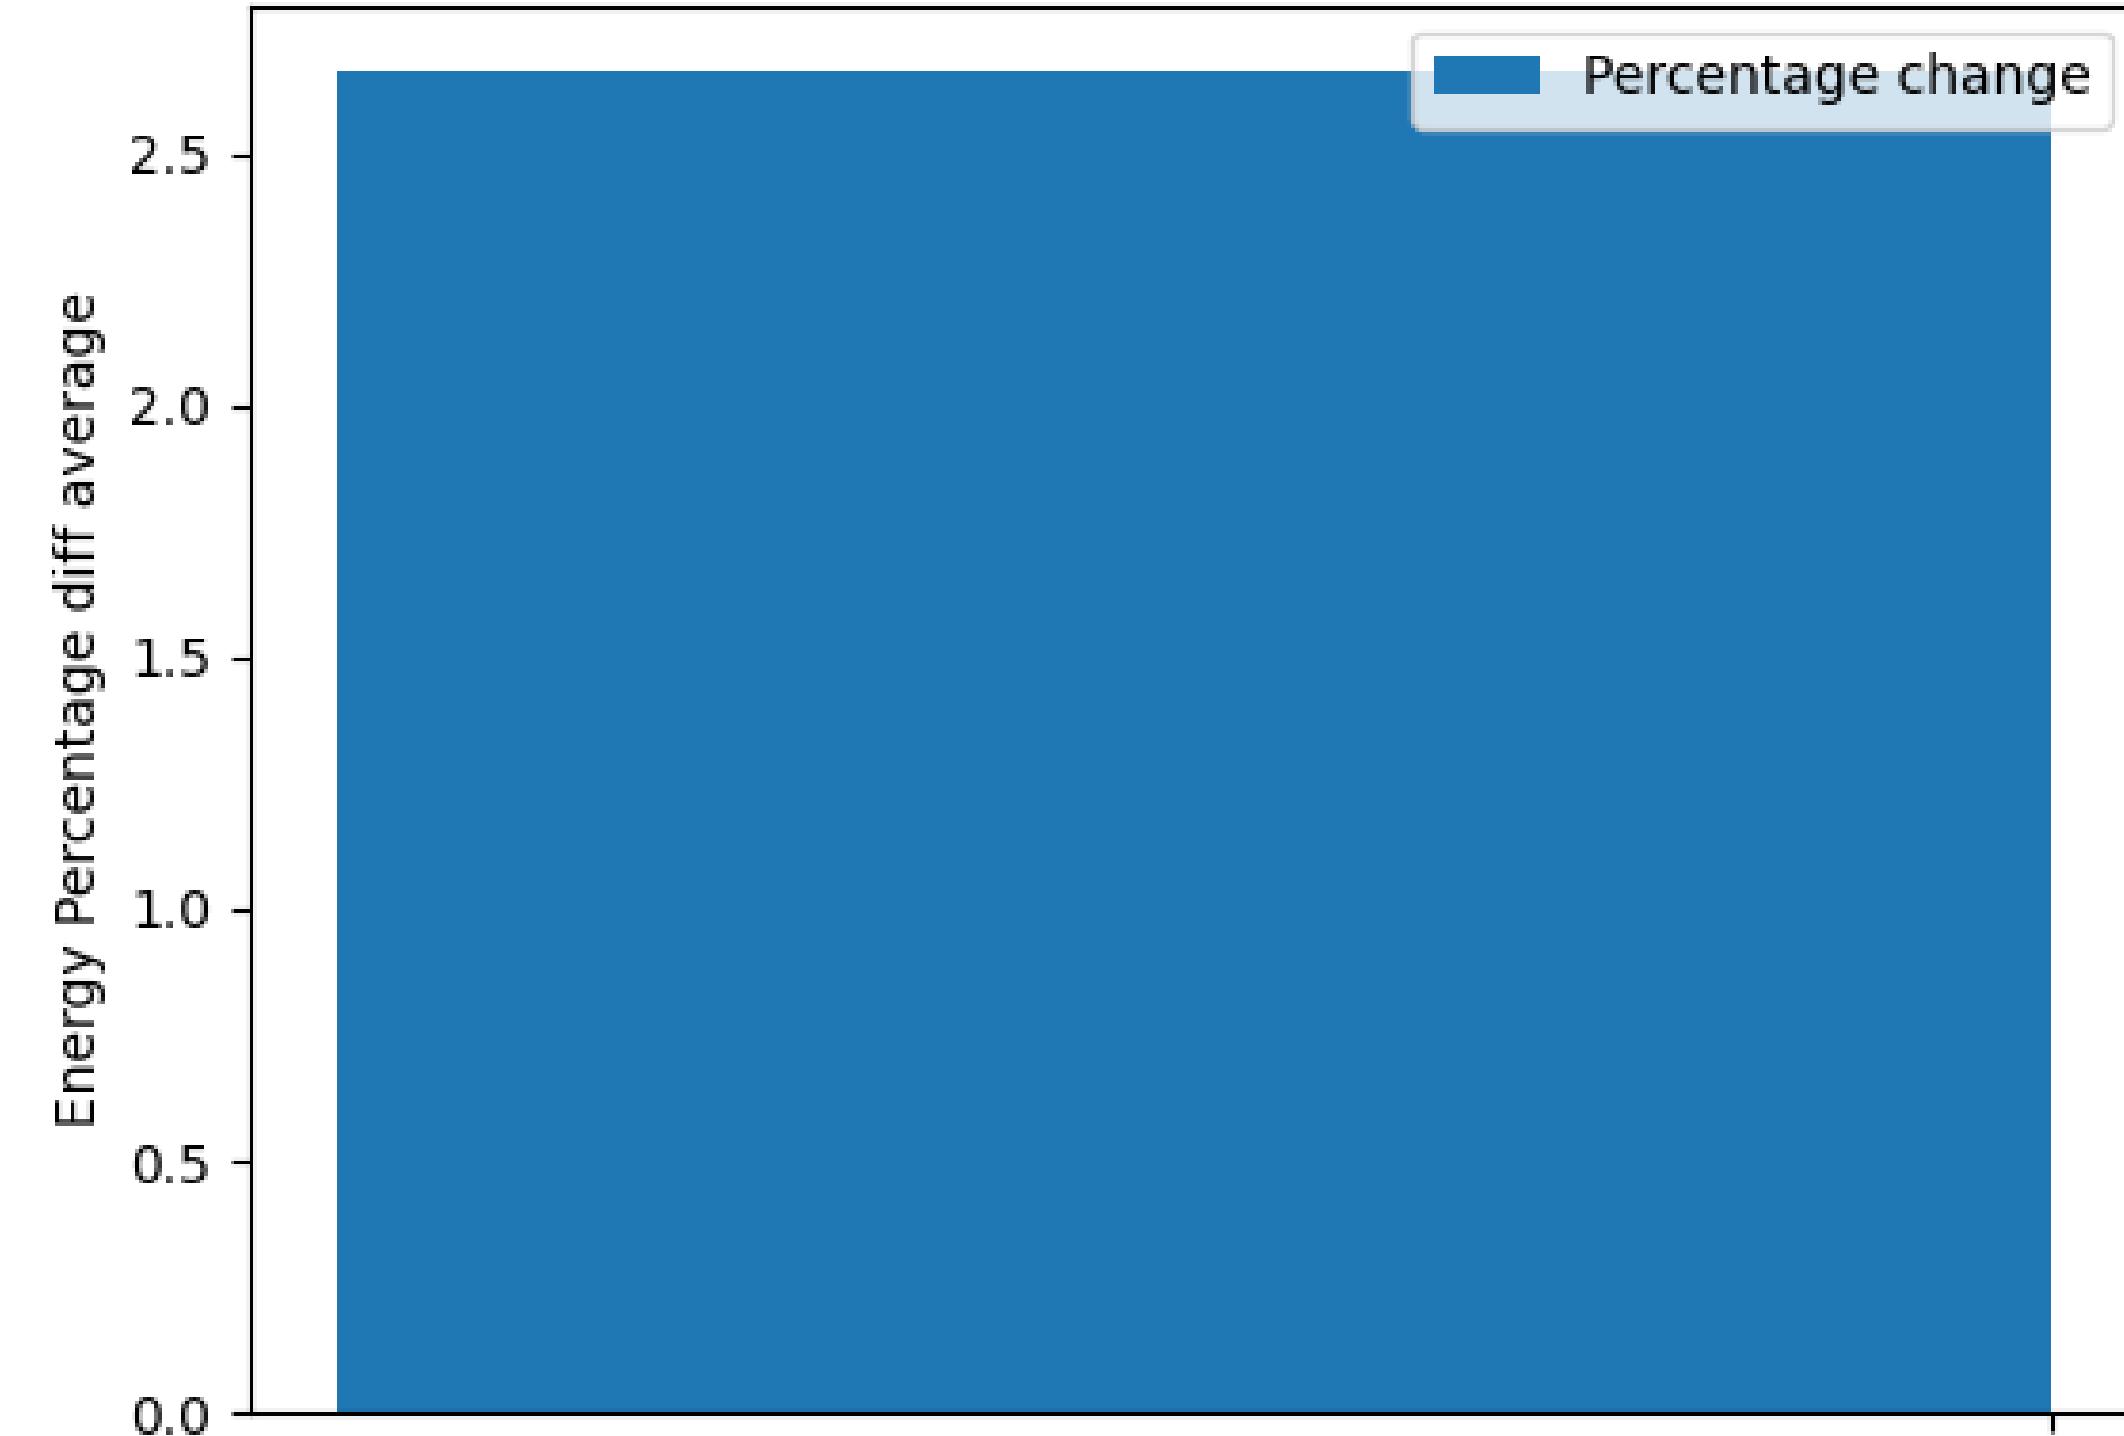
Evaluation Method

Compared actual energy of 24 benchmarks
vs
Estimated energy based on our dataset



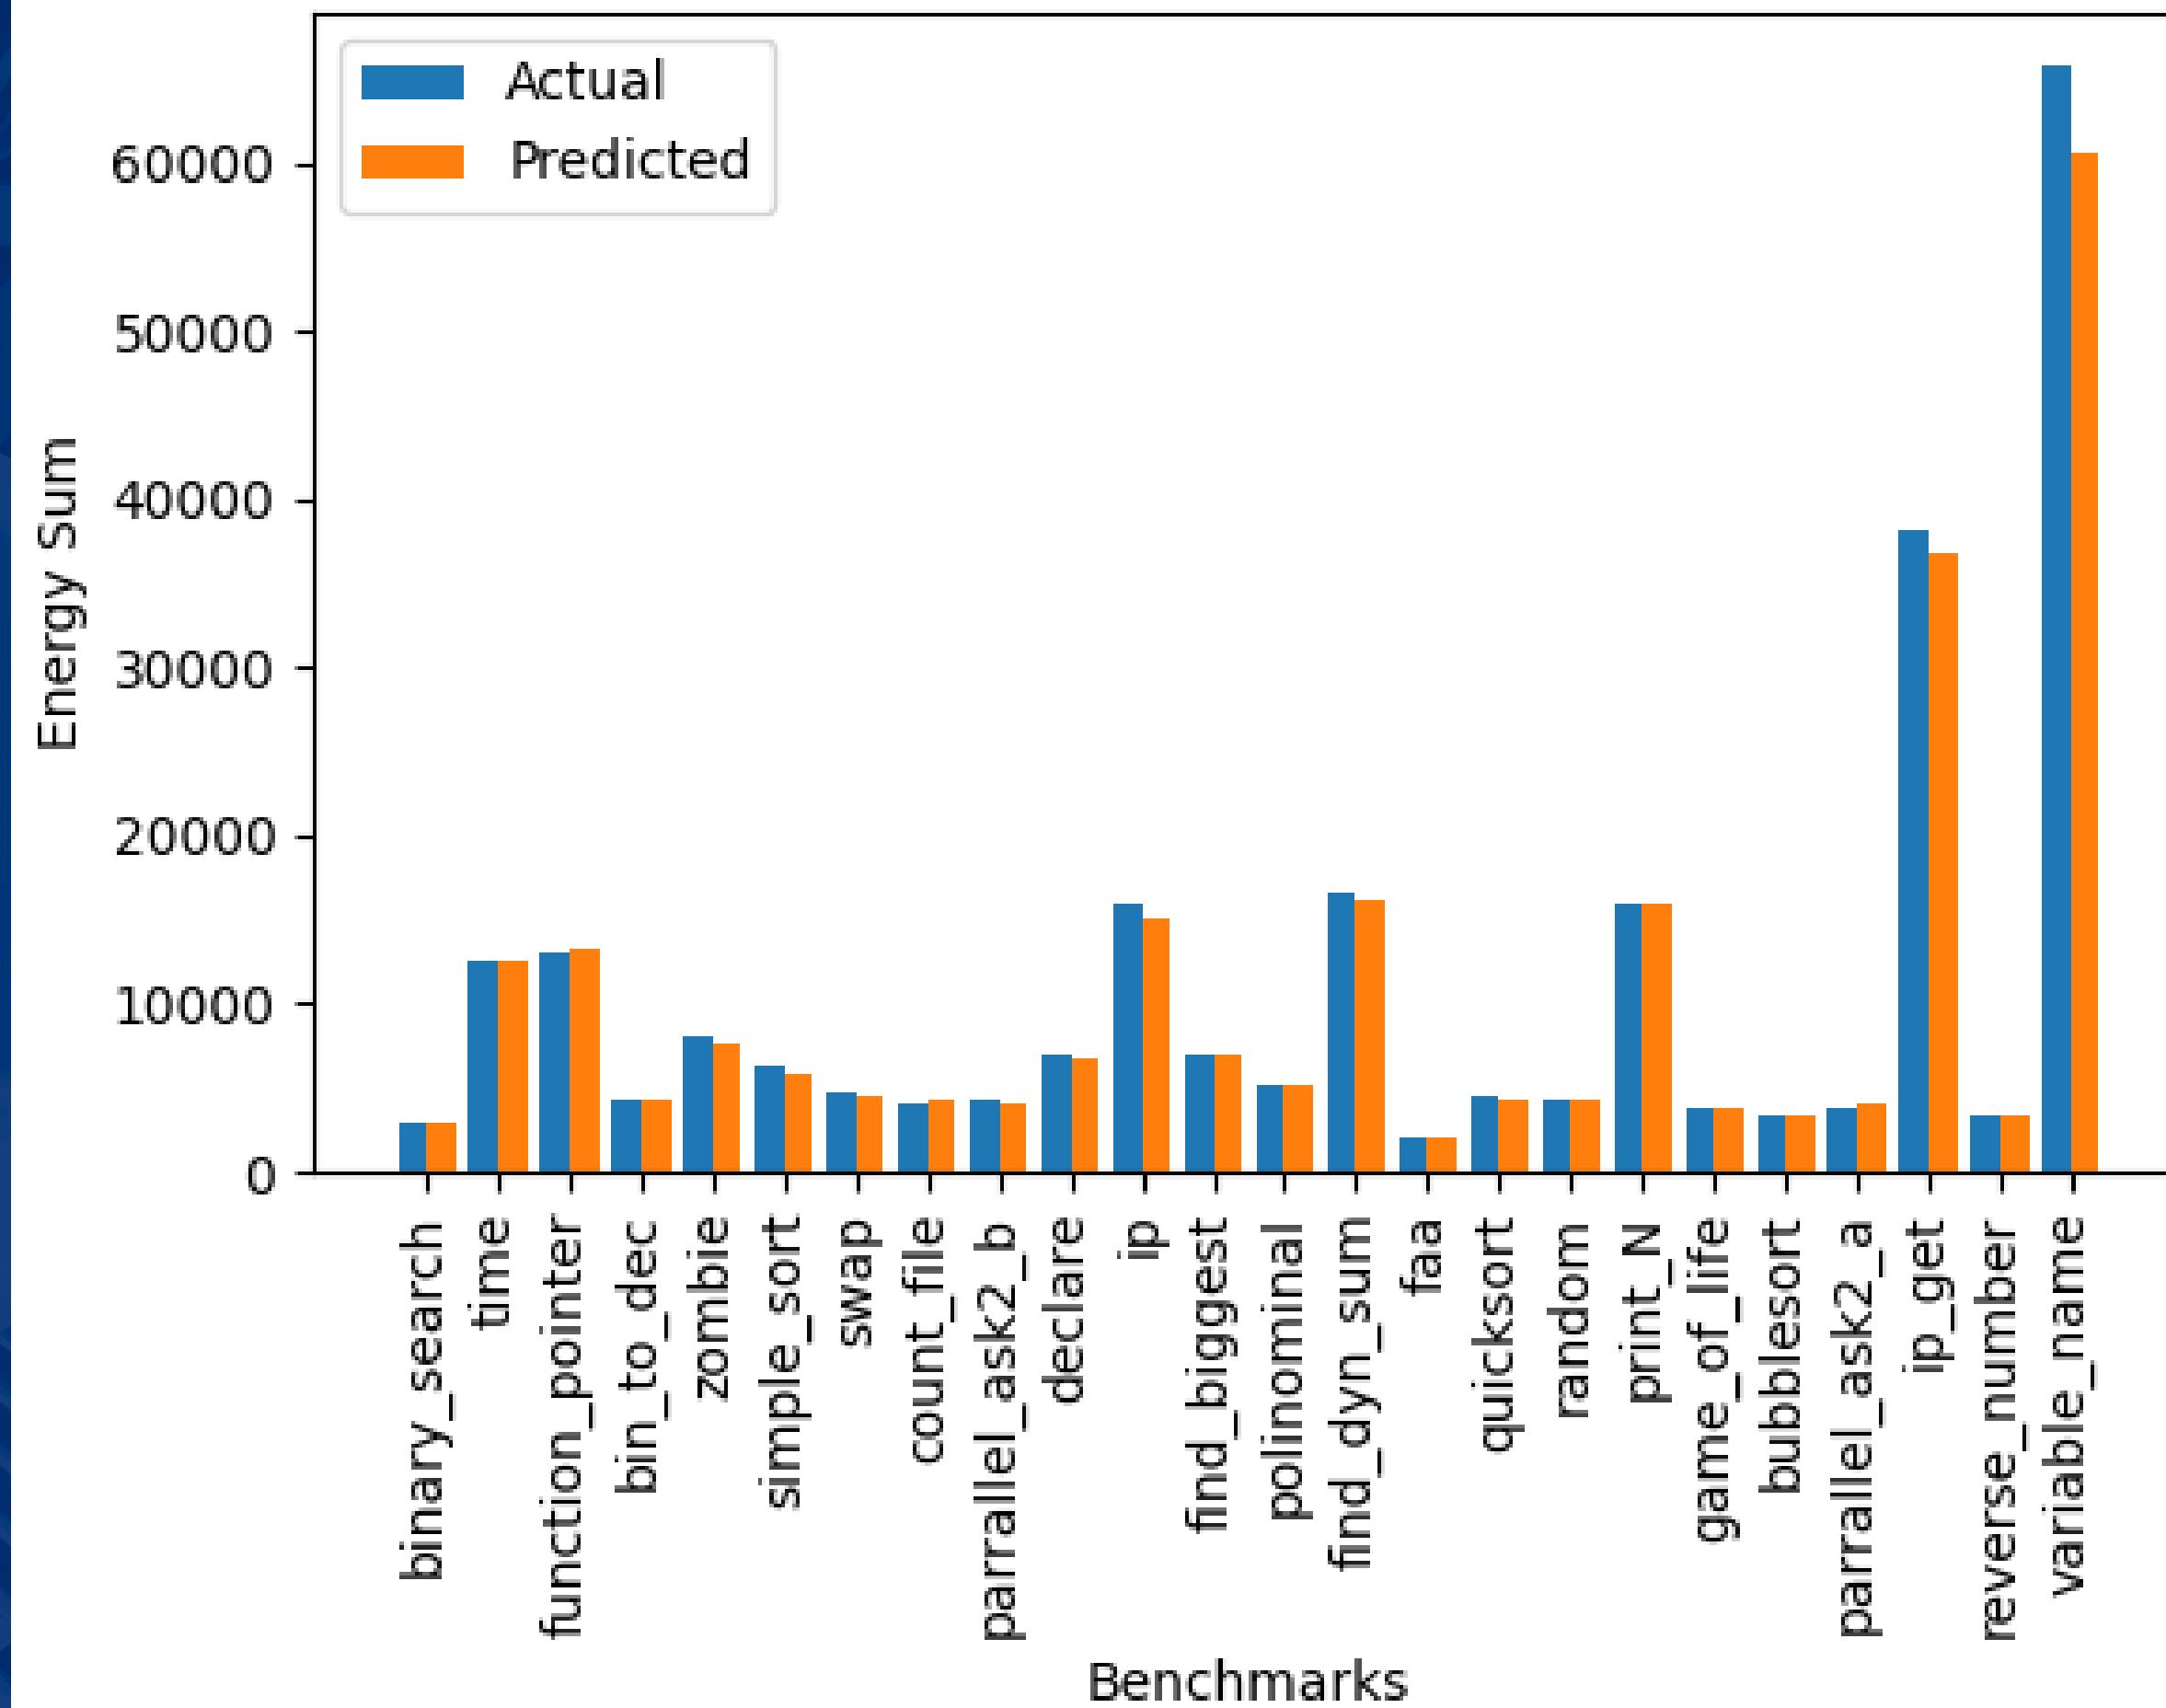
Results

Average difference percentage of Actual energy vs Predicted energy



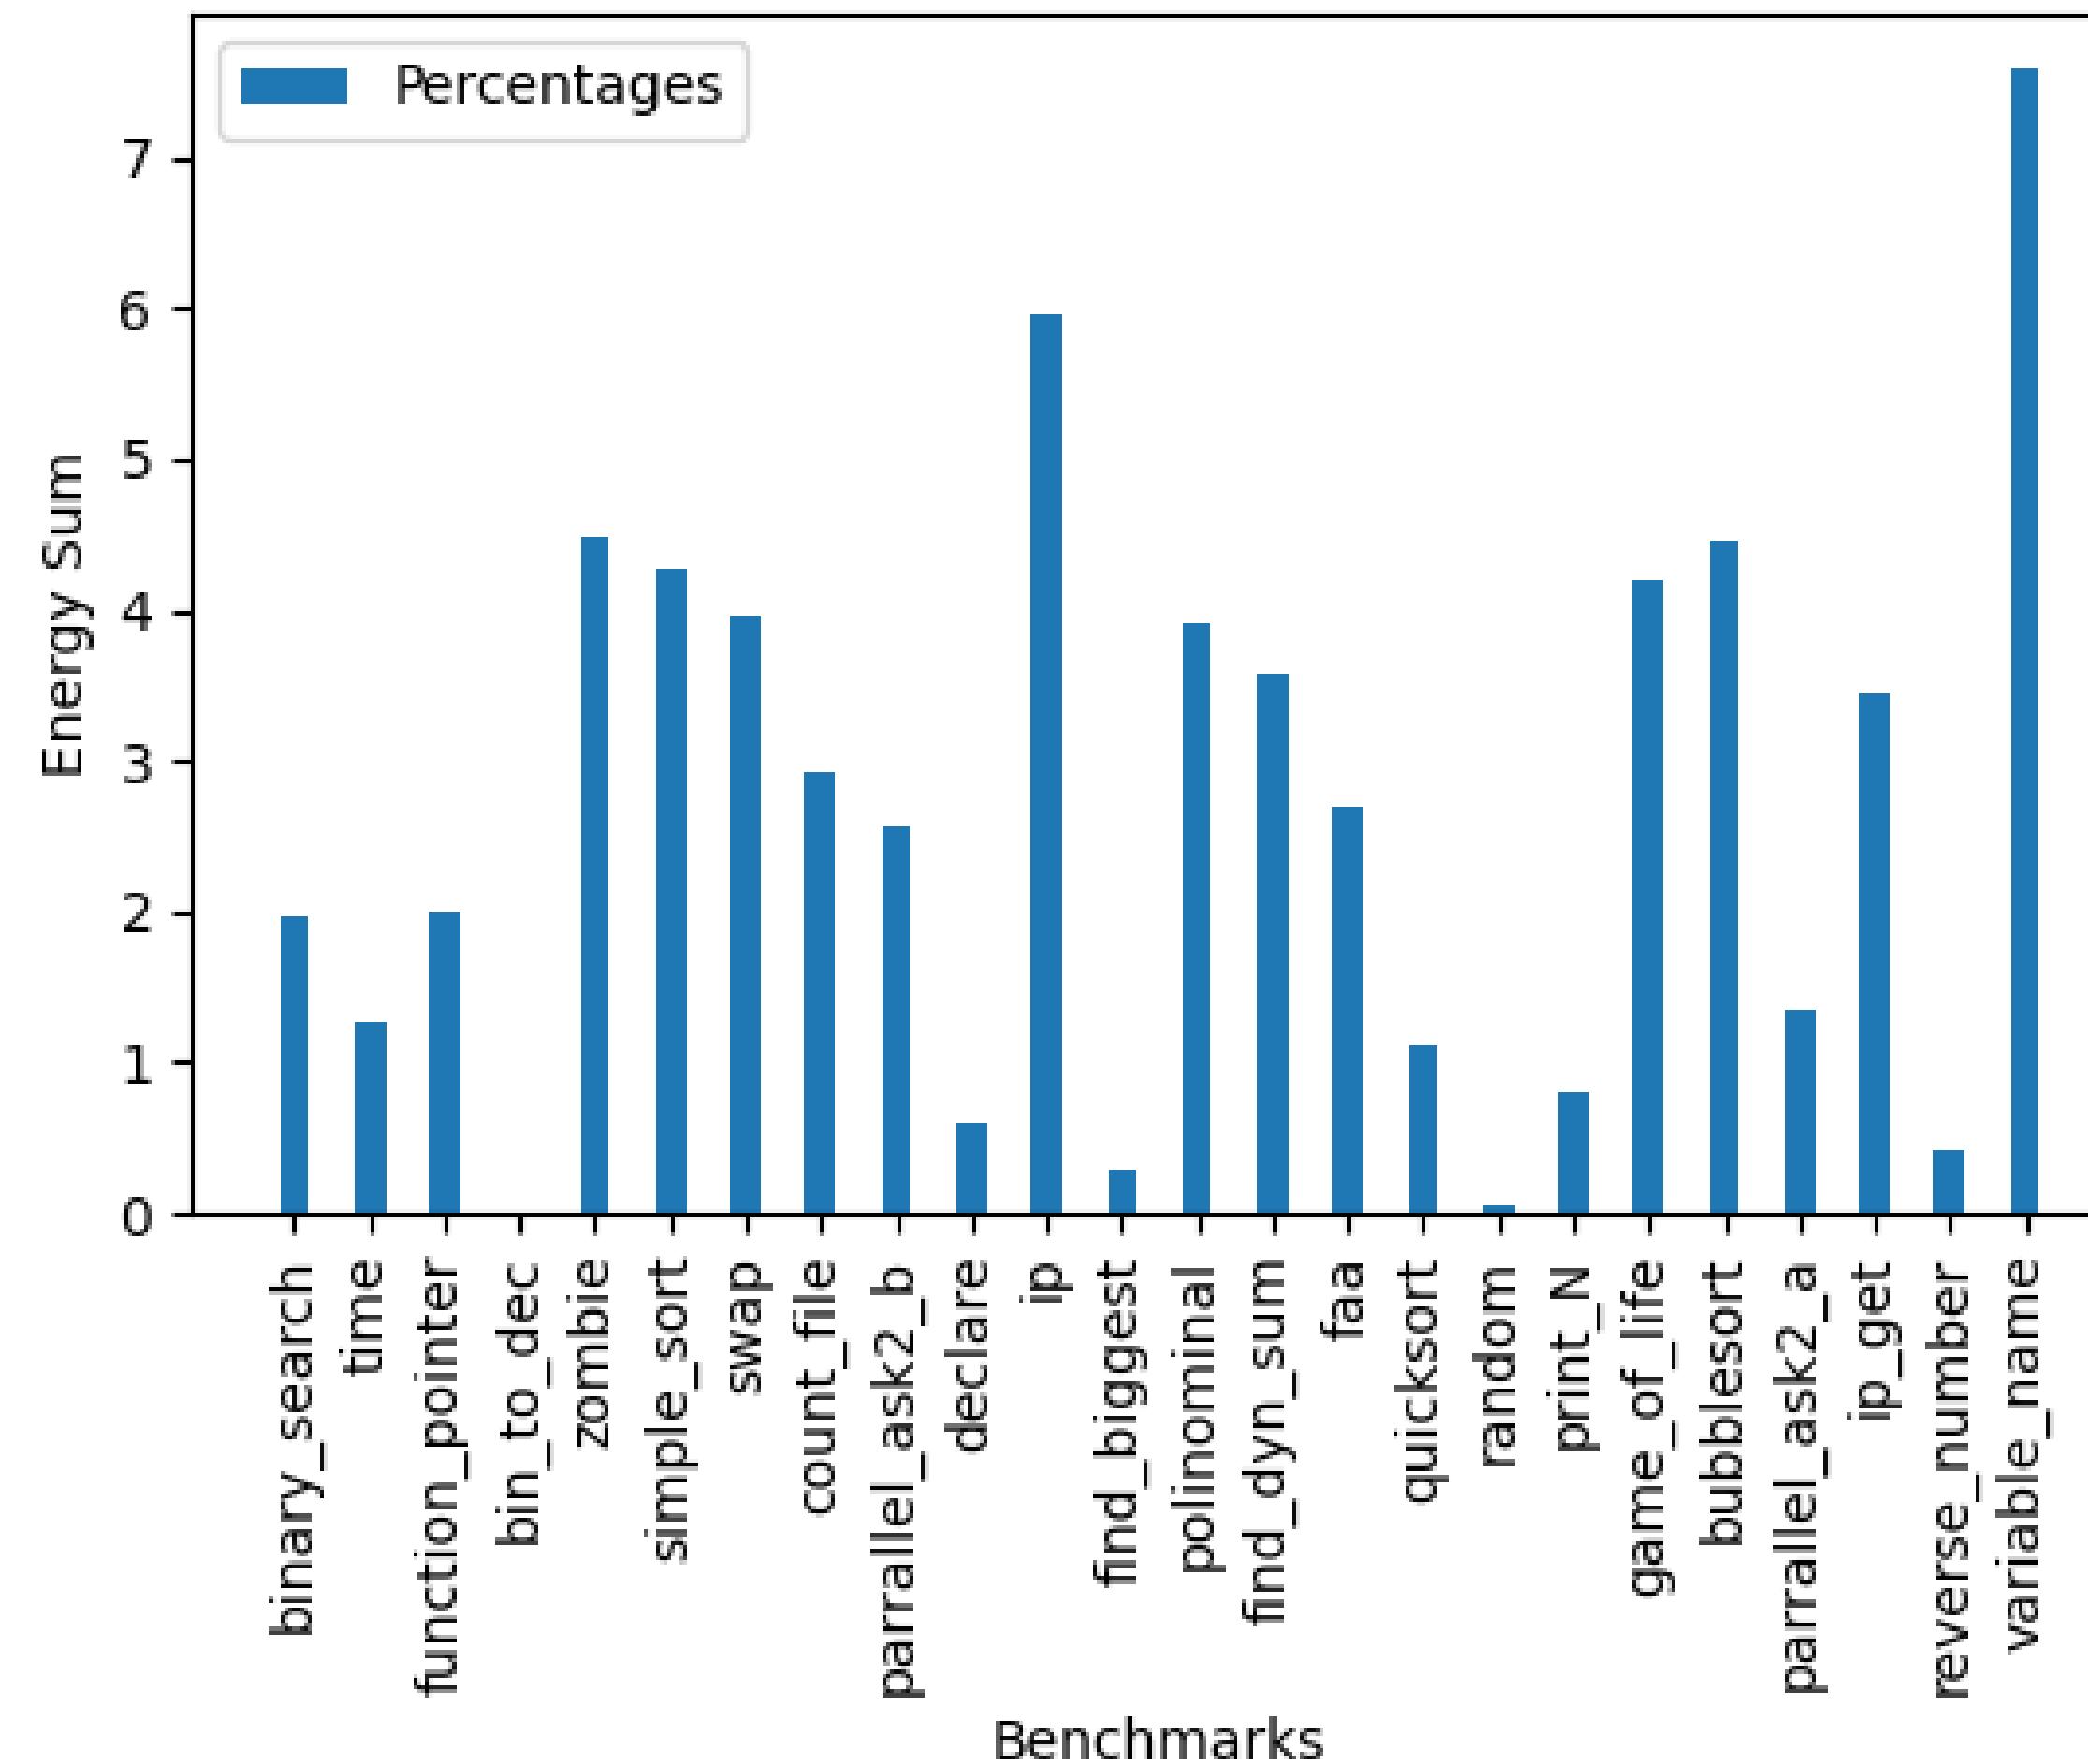


Actual energy sum vs Predicted energy sum





Percentage difference of predicted vs Actual





Comparison with similar research

- Comparable error to the state of the art ,ALEA
- Better accuracy than works with similar granularity
- Better granularity than works with similar accuracy
- Open source and reproducible



Current continuation : Energy dataset for memory and GPU

change
location of
RAPL read
register



repeat
process for
all
benchmarks



create new
dataset for
memory,
gpu, or
uncore



Current continuation : Energy predicting neural network





Future work : Energy reducing code transformations

Software energy reduction techniques
based on code transformation

find energy
costy BBs



replace
them with
others



Conclusions

- Produced reliable energy dataset on a basic block level
- Custom re-production for each system recommended, easy using our tool
- dataset + NN → basis for software reduction techniques
- Binary lifting interesting approach needs more research



Contact Us



ell17072@ntua.mail.gr – Bouras



+30 6980247969



Thank You

**Don't forget to check out the project
codebase on github repo :
https://github.com/jimbou/dataset_tool.git**